

Table of Contents

Quick Topic Links.....	1
Part 1: Foundations & Well-Architected Framework.....	14
What is the difference between Scalability and Elasticity?.....	14
Explain the AWS Shared Responsibility Model.....	15
Region vs. Availability Zone (AZ) vs. Edge Location? Answer: A Region is a geographic area. An AZ is one or more discrete data centers within a region with redundant power/networking. An Edge Location is a CloudFront Point of Presence (PoP) used for caching content closer to users	15
What is Stateful vs. Stateless architecture?.....	15
What are AWS Organizations and Service Control Policies (SCPs)?.....	15
When would you use AWS Control Tower?.....	15
What are the hybrid connectivity options on AWS?.....	15
Part 2: Networking & VPC.....	16
Walk me through the components of a VPC.....	16
What is the difference between a Public and Private Subnet?.....	16
Answer:Public Subnet: Route table has `0.0.0.0/0` pointing to an Internet Gateway. Instances can have public IPs.Private Subnet: Route table has `0.0.0.0/0` pointing to a NAT Gateway.....	16
Security Group vs. NACL?.....	16
VPC Peering vs. Transit Gateway?.....	16
What is a VPC Endpoint?.....	16
NAT Gateway vs. NAT Instance?.....	16
ALB vs. NLB vs. GLB?.....	17
Part 3: Compute & Scaling.....	17
On-Demand vs. Reserved vs. Savings Plans vs. Spot?.....	17
Launch Template vs. Launch Configuration?.....	17
EC2 vs. ECS vs. EKS vs. Fargate vs. Lambda?.....	17
ECS Fargate vs. ECS on EC2?.....	17
What is a Lambda Cold Start and how do you mitigate it?.....	17
Part 4: Storage.....	17
How do you secure an S3 bucket?.....	18
EBS Volume Types?.....	18
EBS vs. EFS vs. FSx? Answer: EBS: Block storage, single AZ, attached to one instance (usually). EFS: NFS file system, multi-AZ, shared by many instances. FSx: Managed Windows/Lustre/NetApp file systems	18
How do you migrate large data to AWS?.....	18
S3 Cross-Region Replication (CRR) vs. Same-Region Replication (SRR)?.....	18
Part 5: Database.....	18
RDS Multi-AZ vs. Read Replicas?.....	18
When to use DynamoDB?.....	18
DynamoDB Partition Key Design?.....	18
ElastiCache Redis vs. Memcached?.....	18
When to use Redshift?.....	18
Part 6: Security & IAM.....	19
What is the Principle of Least Privilege?.....	19
KMS vs. Secrets Manager vs. Parameter Store?.....	19
What is AWS WAF?.....	19
GuardDuty vs. Security Hub vs. Inspector?.....	19
How to secure cross-account access?.....	19
Part 7: Integration & Serverless.....	19
SQS Standard vs. FIFO?.....	19
Lambda Concurrency: Reserved vs. Provisioned?.....	20
Step Functions vs. SQS-driven Lambda?.....	20
Part 8: Design Scenarios (Senior Level).....	20
Design a serverless REST API for 100K RPS.....	20

Design a Data Lake.....	20
Troubleshooting sudden slowness.....	20
Migrating a monolith to AWS.....	20
"Tell me about a system you designed.".....	20
Design an enterprise Agentic AI/RAG platform on AWS for a financial institution.....	21
Edge & Ingress Security.....	21
Core Orchestration Layer (ECS/EKS Private Subnet).....	21
Execution & Data Integration.....	21
Inference & HITL Gateways.....	21
Part 1: Amazon Bedrock & Foundation Models (FM).....	22
What are the key features of Amazon Bedrock Agents?.....	22
Explain the concept of "Model Distillation" in Bedrock.....	22
How does Amazon Bedrock handle data privacy?.....	23
What is the difference between On-Demand and Provisioned Throughput in Bedrock?.....	23
What are "Guardrails" in Amazon Bedrock?.....	23
How do you evaluate FM performance in Bedrock?.....	23
Part 2: RAG (Retrieval-Augmented Generation).....	23
What vector databases are supported natively by Bedrock Knowledge Bases?.....	23
How does Bedrock handle "Chunking" in RAG? Answer: Bedrock Knowledge Bases offer automatic chunking strategies (fixed-size, hierarchical, or semantic) during data ingestion. You can also define custom chunking logic using Lambda functions if the default strategies don't fit your document structure	23
What is the role of Amazon Titan Embeddings in RAG?.....	23
How do you improve retrieval accuracy in a RAG system?.....	24
Can Bedrock Knowledge Bases connect to real-time data?.....	24
Part 3: Agentic AI & Amazon Q.....	24
How does Amazon Q Business differ from Q Developer?.....	24
What is "Agentic AI" in the context of AWS?.....	24
How do you build a custom agent in Bedrock?.....	24
How does Amazon Q handle enterprise security?.....	24
What are "Q Apps"?.....	24
Part 4: Deployment, Monitoring, Logging & Evaluation.....	25
How do you monitor Bedrock API usage and costs?.....	25
How do you log prompts and responses for auditing?.....	25
What is Amazon Bedrock Flows?.....	25
How do you test for "Hallucinations" in a RAG app?.....	25
What is the role of AWS Lambda in GenAI architectures?.....	25
How do you secure GenAI applications against Prompt Injection?.....	25
What is "Human-in-the-Loop" (HITL) in Bedrock?.....	25
How do you manage versioning for Bedrock Agents?.....	25
What is the "Well-Architected Lens for GenAI"?.....	26
How do you handle rate limiting in Bedrock?.....	26
What is the role of Amazon EventBridge in Agentic AI?.....	26
Part 1: EKS Fundamentals & Architecture.....	26
Explain the EKS Control Plane vs. Data Plane.....	26
What is the difference between EKS on EC2 and EKS on Fargate?.....	26
What is an EKS Add-on?.....	26
How does EKS handle version upgrades?.....	27
What is EKS Auto Mode? Answer: A newer feature (2024-2026) that automates node provisioning, scaling, and optimization. It automatically selects the best instance types, manages spot/on-demand mixing, and handles OS patching, reducing the need for manual Node Group configuration	27
Part 2: Networking & Security in EKS.....	27
What is the difference between Security Groups for Pods and Node Security Groups?.....	27
How do you secure access to the EKS API Server?.....	27
What is IRSA (IAM Roles for Service Accounts)?.....	27
How do you expose an EKS application to the internet?.....	27
What is the AWS Load Balancer Controller?.....	27
Part 3: Storage & Stateful Workloads How do you provide persistent storage to EKS Pods?	

.....	28
What is the difference between EBS and EFS in EKS?.....	28
How do you handle secrets in EKS?.....	28
Part 4: EKS for AI/GenAI & ML Workloads.....	28
How do you manage GPUs in EKS?.....	28
What is Karpenter and why is it preferred over Cluster Autoscaler?.....	28
How do you optimize costs for AI inference on EKS?.....	28
Can EKS integrate with Amazon Bedrock?.....	28
What is KServe or Seldon Core in EKS?.....	29
Part 5: Monitoring, Logging & Observability.....	29
How do you centralize logs from EKS?.....	29
What is X-Ray and how does it help in EKS?.....	29
How do you debug a "CrashLoopBackOff" in EKS?.....	29
Part 6: Advanced EKS & GitOps.....	29
What is EKS Anywhere?.....	29
How do you handle multi-tenancy in EKS?.....	29
What is the role of AWS App Mesh in EKS?.....	30
How do you secure the supply chain in EKS?.....	30
Case Scenario.....	30
Phase 1: Establish the Facade (The "Vine").....	31
Phase 2: Build the New Microservice (The "New Growth").....	31
Phase 4: Traffic Shifting (The "Strangling").....	31
What is the role of AWS Application Migration Service (MGN)?.....	32
How do you handle database modernization from Oracle/SQL Server to Aurora?.....	32
What are the biggest cultural challenges in legacy modernization?.....	32
Part 2: Developer Experience (DevEx) & Internal Developer Platforms (IDP).....	32
How does AWS Proton fit into Developer Transformation?.....	32
What metrics do you use to measure Developer Productivity?.....	32
How do you implement "Golden Paths" or "Paved Roads"?.....	33
What is the role of Backstage in AWS environments?.....	33
Part 3: CI/CD at Enterprise Scale.....	33
How do you implement Blue/Green Deployments on AWS?.....	33
What is GitOps and how do you enforce it in a large organization?.....	33
How do you secure the CI/CD pipeline?.....	33
What is the "Shift Left" security model?.....	33
Part 4: Strategic Consulting & Leadership Principles.....	34
Describe a time you influenced a technical decision against a customer's initial preference.	34
How do you assess an organization's "Cloud Maturity"?.....	34
What is the role of a "Center of Excellence" (CoE)?.....	34
How do you align technical transformation with business goals?.....	34
Part 5: Advanced Modernization Patterns.....	34
How do you handle distributed transactions in a microservices architecture?.....	34
What is the "Anti-Corruption Layer" (ACL)?.....	35
How do you modernize .NET Framework applications?.....	35
What is the role of Amazon MQ in modernization?.....	35
Part 6: AI/GenAI in Developer Transformation.....	35
How do you govern GenAI usage in an enterprise?.....	35
What is "Agentic CI/CD"?.....	35
How do you measure the ROI of GenAI in development?.....	35
What is the biggest risk of adopting GenAI in coding?.....	35
Explain about aws migration services and how can it help to on premises or other cloud services to aws.....	36
Discovery & Planning: "What do we have?".....	36
Data Transfer: "How do we get the bits there?".....	36
Server & VM Migration (Rehost / Lift-and-Shift).....	36
Database Migration (Replatform / Refactor).....	37
Mainframe & Legacy Modernization (Refactor / Replatform).....	37
Application & Code Modernization (Refactor).....	37

End-of-Life & Decommissioning (Retire).....	37
Strategic Framework: The 6 Rs in Practice.....	38


For Amazon Senior Solutions Architect interview, these are exactly the questions I would prepare. Don't memorize long answers; memorize the architecture thinking and a 30–60 second answer for each.

How can you make your application scalable for a big traffic day?

I would answer:

“I would first identify the expected traffic, TPS, latency SLA and read/write ratio. Then I would remove single points of failure and make the application horizontally scalable.”

AWS architecture:



Read Replicas Key techniques:

- CloudFront → absorb/cache static traffic
- WAF → protect application
- ALB → distribute traffic
- EC2 Auto Scaling / ECS / EKS → horizontal scaling
- SQS → decouple burst traffic
- ElastiCache → reduce database load
- RDS Multi-AZ → availability
- Read replicas → read scaling
- DynamoDB → extremely high-scale key/value workloads

S3 → virtually unlimited object storage

CloudWatch → monitor and automatically scale The important Senior Architect statement:

“I would design for elasticity rather than simply provisioning a larger server.”

How do you achieve DR for your cloud application?


Start with:

“I would first define RPO and RTO because the DR architecture depends on the business requirements.”

RPO = How much data can I afford to lose?

RTO = How quickly must I recover? Then:

Primary Region



DR Region Strategies:

Backup & Restore



Pilot Light



Warm Standby



Active/Active For example:

Route 53



Region A Production



| replication

▼ Region B DR

AWS services:

Route 53

AWS Backup

S3 Cross-Region Replication


Aurora Global Database / appropriate database replication

DynamoDB Global Tables where appropriate
CloudFormation/CDK/Terraform
CloudWatch
Systems Manager Strong interview statement:


“DR isn't simply having a backup. I need to demonstrate that I can restore the application, data, configuration and dependencies within the agreed RTO.”

How do you secure your application on the cloud?

Use the AWS Well-Architected security mindset: SECURITY



Architecture:



Private Database Key points:

Network

VPC
public/private subnets
Security Groups
NACLs
VPC endpoints
no public database

Identity

IAM
IAM Roles
least privilege
MFA
temporary credentials
Cognito/OIDC for application users

Data

KMS
encryption at rest
TLS in transit
Secrets Manager
S3 bucket policies

Monitoring

CloudTrail
GuardDuty
Security Hub
Config
CloudWatch Excellent answer:

“I use defense in depth. I don't depend on a single security control.”


Describe an architecture you designed


For you, I would use your AI/GenAI enterprise experience rather than inventing a generic ecommerce system.

You can say:

“One architecture I worked on was an enterprise AI platform integrating RAG, agentic workflows and risk technology.”

Then draw:





↓
Bedrock / LLM

↓
Guardrails

↓
Response Ingestion:

Documents

↓
S3

↓
EventBridge

↓
SQS

↓
Processing

↓
OCR / PII / Classification

↓
Chunking

↓
Embeddings

↓
OpenSearch

Then explain the trade-offs:

security

latency

scalability

model cost

RAG quality

observability

multi-tenancy

This is much stronger than simply saying:

“I built an application using EC2 and RDS.”

Biggest challenge during designing your application on cloud?

Don't say:

“The biggest challenge was configuring AWS.” Instead use an architectural challenge.

For your background, I would use:

“The biggest challenge was balancing AI quality, latency, cost and enterprise security.”

Then explain:

Higher quality model

↓

Higher cost + latency

More agents

↓

Better specialization

↓

More LLM calls

↓

Higher cost + latency

More security controls

↓

Potential additional latency Then:

“I addressed this by measuring the complete request path, using caching and retrieval optimization, routing simpler tasks to smaller models, limiting agent/tool calls, and applying security controls at the authorization boundary rather than relying on the model.”

That's a very good Senior Architect answer.

How do you pick one service versus another as a Solutions Architect?

This is probably one of the most important questions. Never answer:

“I choose the service I know best.”

Say:

“I start with requirements rather than services.” Use this framework:

Business Requirements

↓

Functional Requirements

↓

Non-functional Requirements

↓

Constraints



Architecture Options



Trade-off Analysis



Service Selection Evaluate:

Scalability

Availability

Performance

Security

Cost

Operational complexity

Integration

Team skills

Compliance

Vendor/service maturity Example:

“If I need a relational transactional database with complex joins, I might choose Aurora/RDS. If I need massive horizontal key-value scaling with predictable access patterns, DynamoDB may be better. I wouldn't choose DynamoDB simply because it scales.”

Excellent phrase:

“The requirement drives the service; the service doesn't drive the requirement.”

Difference between SQL and NoSQL?

SQL

Structured schema Tables

Rows Relationships SQL

ACID transactions Complex queries

Examples:

Aurora

RDS PostgreSQL

RDS MySQL

Good for:

Banking Orders Payments

Financial transactions ERP

NoSQL

Flexible schema Horizontal scaling Access-pattern driven High throughput


AWS examples:

DynamoDB
DocumentDB
Neptune

Keyspaces Example:

SQL

Customer



Orders

Payments versus DynamoDB:

PK = CUSTOMER#123 SK = ORDER#456

The key interview point:

“NoSQL isn't simply SQL without a schema. The data model and access patterns are fundamentally different.”

What is cloud computing?

Short answer:

“Cloud computing is the on-demand delivery of computing resources such as compute, storage, databases, networking and AI services over the internet, with elastic capacity and usage-based economics.”

Then mention the key benefits: On-demand

Elastic

Global

Pay-as-you-go Highly available Managed services Automation

Traditional:

Buy servers



Install



Configure



Maintain



Capacity planning

Cloud: Requirement



Provision API

↓

Scale automatically

↓

Pay for usage

What is AWS Service X?

This question tests whether you understand services rather than memorized definitions. Use this template:

AWS Service X

What problem does it solve?

How does it work?

When would I use it?

When would I NOT use it?

What are the alternatives?

What are the trade-offs?

For example, if they ask:

“What is SQS?”

Don't just say:

“SQS is a queue service.”

Say:

“Amazon SQS is a managed message queuing service used to decouple distributed components. I would use it when producers and consumers should operate independently, particularly when I need to absorb traffic spikes, handle asynchronous processing or improve resilience.”

Then:

Producer

↓

SQS

↓

Consumer Benefits:

decoupling

buffering

retry

asynchronous processing

elasticity

The 9 questions — your mental cheat sheet

Before the interview, remember this: SCALABILITY

→ Horizontal scaling + caching + queues + autoscaling

DR

→ RPO + RTO + backups + replication + tested recovery

SECURITY

→ IAM + VPC + encryption + WAF + monitoring

ARCHITECTURE

→ Requirements → design → trade-offs → implementation

CHALLENGE

→ Quality vs cost vs latency vs security

SERVICE SELECTION

→ Requirements → trade-offs → service

SQL vs NoSQL

→ Transactions/relationships vs scale/access patterns

CLOUD

→ On-demand + elastic + managed + pay-as-you-go

AWS SERVICE X

→ Problem → usage → alternative → trade-off

For your interview, the biggest thing to practice is not AWS service definitions. It is being able to say:

“Given these requirements, I would choose X instead of Y because...” That is what separates a cloud engineer from a Solutions Architect.

The interview landscape in 2026 focuses heavily on trade-offs, cost optimization, and design judgment rather than just memorizing service features .

Part 1: Foundations & Well-Architected Framework

What are the six pillars of the AWS Well-Architected Framework?

Answer: Operational Excellence, Security, Reliability, Performance Efficiency, Cost Optimization, and Sustainability. Senior architects should note that these pillars often trade off against each other (e.g., multi-AZ improves Reliability but increases Cost) .

What is the difference between Scalability and Elasticity?

Answer: Scalability is the ability to handle increased load (vertical or horizontal). Elasticity is the ability to *automatically scale up and down based on demand. AWS is designed to be elastic, whereas on-premise systems are typically only scalable .

Explain the AWS Shared Responsibility Model.

Answer: AWS is responsible for security of the cloud (physical infrastructure, hypervisor). The customer is responsible for security in the cloud (IAM, network configuration, OS patching for EC2, data encryption). The boundary shifts depending on the service (e.g., less customer responsibility in Lambda vs. EC2) .

Region vs. Availability Zone (AZ) vs. Edge Location?

Answer: A Region is a geographic area. An AZ is one or more discrete data centers within a region with redundant power/networking. An Edge Location is a CloudFront Point of Presence (PoP) used for caching content closer to users .

What is Stateful vs. Stateless architecture?

Answer: Stateless: No local session data; any instance can serve any request (state is externalized to DB/Cache). Scales horizontally easily. Stateful: Instances hold local state (e.g., in-memory sessions). Requires session affinity or complex replication. Modern AWS designs prefer stateless tiers .

What are AWS Organizations and Service Control Policies (SCPs)?

Answer: Organizations allow centralized management of multiple accounts. SCPs are guardrails applied at the OU or account level to define the maximum permission boundary. SCPs cannot grant permissions; they can only restrict what IAM otherwise allows .


When would you use AWS Control Tower?

Answer: Use Control Tower to automate the setup of a secure, multi-account landing zone. It provides baseline guardrails, centralized logging, and identity federation. It is recommended when managing more than 3-4 AWS accounts .

What are the hybrid connectivity options on AWS?

Answer: Site-to-Site VPN (IPsec over internet, quick setup), AWS Direct Connect (dedicated fiber, predictable performance), and Transit Gateway (hub for connecting VPCs and on-prem networks). Most enterprises use Direct Connect with VPN as a backup .

Part 2: Networking & VPC



Walk me through the components of a VPC.

Answer: CIDR block, Subnets (public/private per AZ), Route Tables, Internet Gateway (IGW), NAT Gateway, Security Groups (instance-level), NACLs (subnet-level), and VPC Endpoints .

What is the difference between a Public and Private Subnet?

Answer: Public Subnet: Route table has `0.0.0.0/0` pointing to an Internet Gateway. Instances can have public IPs. Private Subnet: Route table has `0.0.0.0/0` pointing to a NAT Gateway. Instances have no public IP and can only initiate outbound traffic .

Security Group vs. NACL?

Answer: Security Group: Stateful, instance-level, allow rules only. NACL: Stateless, subnet-level, allow and deny rules. Security groups are evaluated first .

VPC Peering vs. Transit Gateway?

Answer: VPC Peering: Point-to-point, non-transitive (A-B and B-C does not mean A-C). Transit Gateway: Hub-and-spoke model, supports transitive routing, and connects many VPCs, VPNs, and Direct Connect gateways. Use TGW for complex networks .

What is a VPC Endpoint?

Answer: Allows private connectivity to AWS services without traversing the internet. Gateway Endpoints (free) for S3/DynamoDB. Interface Endpoints (paid, PrivateLink) for most other services .

NAT Gateway vs. NAT Instance?

Answer: NAT Gateway: Managed, highly available within an AZ, scales to 45 Gbps. NAT Instance: Self-managed EC2, cheaper at low volume but requires manual HA and scaling. Default to NAT Gateway for production .

ALB vs. NLB vs. GLB?

Answer: ALB: Layer 7 (HTTP/HTTPS), path-based routing. NLB: Layer 4 (TCP/UDP), ultra-high performance, static IP. Gateway Load Balancer (GLB): Used for inserting third-party virtual appliances (firewalls/IDS) into the network path .

Part 3: Compute & Scaling

How do you choose EC2 instance types?

Answer: M-series: General purpose. C-series: Compute optimized. R/X-series: Memory optimized. I/D-series: Storage optimized. P/G/Inf-series: GPU/Accelerated computing. Use Graviton (ARM) instances for better price/performance where supported .

On-Demand vs. Reserved vs. Savings Plans vs. Spot?

Answer: On-Demand: Pay-as-you-go. Reserved/Savings Plans: 1-3 year commitment for discounts. Spot: Up to 90% discount but can be interrupted with 2-minute notice. Mix these for cost optimization .

What scaling policies exist for Auto Scaling Groups?

Answer: Target Tracking (keep CPU at X%), Step Scaling, Scheduled Scaling, and Predictive Scaling (ML-driven). Combine target tracking with scheduled scaling for predictable spikes .

Launch Template vs. Launch Configuration?

Answer: Launch Templates are the modern standard. They support versioning, mixed instance types, and Spot allocation strategies. Launch Configurations are deprecated .

EC2 vs. ECS vs. EKS vs. Fargate vs. Lambda?

Answer: EC2: Full control. ECS: AWS-native container orchestration. EKS: Managed Kubernetes. Fargate: Serverless containers (no node management). Lambda: Event-driven functions. Choose Lambda for short events, Fargate for containers, EKS for K8s ecosystem, EC2 for legacy .

ECS Fargate vs. ECS on EC2?

Answer: Fargate: No server management, pay per task. EC2: Cheaper at sustained high scale, supports GPUs/custom AMIs. Default to Fargate for new workloads unless specific hardware needs exist .

What is a Lambda Cold Start and how do you mitigate it?

Answer: Latency during the first invocation after inactivity. Mitigate using Provisioned Concurrency, smaller deployment packages, faster runtimes (Go/Rust), or avoiding VPC attachment if not needed .

Part 4: Storage

S3 Storage Classes – Pick one for each scenario.

Answer: Standard: Frequent access. Intelligent-Tiering: Unknown patterns. Standard-IA: Infrequent access. Glacier Instant Retrieval: Archive with millisecond access. Glacier Deep Archive: Lowest cost, 12+ hour retrieval .

How do you secure an S3 bucket?

Answer: Block Public Access, Bucket Policies (least privilege), VPC Endpoints, Encryption (SSE-KMS/S3), Enforce HTTPS, Versioning, and Object Lock for ransomware protection .

EBS Volume Types?

Answer: gp3: General purpose SSD (default, decoupled IOPS). io2: High IOPS for databases.st1/sc1: Throughput/Cold HDD. Default to gp3 for new workloads .

EBS vs. EFS vs. FSx?

Answer: EBS: Block storage, single AZ, attached to one instance (usually). EFS: NFS file system, multi-AZ, shared by many instances. FSx: Managed Windows/Lustre/NetApp file systems .

How do you migrate large data to AWS?

Answer: DataSync: Online sync. Snowball: Physical appliance for TBs/PBs. S3 Transfer Acceleration: For global uploads over internet. Choose based on network bandwidth and data size .

S3 Cross-Region Replication (CRR) vs. Same-Region Replication (SRR)?

Answer: CRR: For DR and compliance. SRR: For log aggregation or cross-account access. Both require versioning enabled .

Part 5: Database**RDS vs. Aurora?**

Answer: RDS: Managed traditional engines (MySQL, Postgres, etc.). Aurora: AWS-built engine compatible with MySQL/Postgres, separated storage layer, 5x performance, automatic 6-way replication across 3 AZs. Aurora is preferred for production .

RDS Multi-AZ vs. Read Replicas?

Answer: Multi-AZ: Synchronous standby for HA/Failover (not for reads). Read Replicas: Asynchronous copies for scaling read throughput. Use both for HA and performance .

When to use DynamoDB?

Answer: Single-digit millisecond latency, key-value/document workloads, unpredictable scale. Avoid complex relational queries or large blobs. Use Global Tables for multi-region active-active .

DynamoDB Partition Key Design?

Answer: Choose a high-cardinality key to avoid "hot partitions." Use Composite Keys (Partition + Sort) for query flexibility. Use GSIs for alternative access patterns .

ElastiCache Redis vs. Memcached?

Answer: Redis: Persistence, pub/sub, complex data structures, Multi-AZ HA. Memcached: Simple, multithreaded, no persistence. Default to Redis .

When to use Redshift?

Answer: Petabyte-scale data warehousing for SQL analytics. Use Redshift Serverless for variable workloads. Use Athena for ad-hoc SQL on S3 .---

Part 6: Security & IAM

IAM User vs. Role vs. Group?

Answer: User: Long-lived credentials (avoid for apps). Role: Temporary credentials via STS (used by services/apps). Group: Collection of users for permission management. Best practice: No IAM users for humans (use Identity Center) or workloads (use Roles) .

What is the Principle of Least Privilege?

Answer: Grant only the permissions required. Use IAM Access Analyzer to find unused permissions. Start broad and tighten based on CloudTrail logs .

KMS vs. Secrets Manager vs. Parameter Store?

Answer: KMS: Manages encryption keys. Secrets Manager: Rotates DB credentials (\$0.40/secret). Parameter Store: Free hierarchical config storage. Use Secrets Manager for rotating secrets, Parameter Store for config .

What is AWS WAF?

Answer: Web Application Firewall. Protects against OWASP Top 10, bots, and rate limiting. Attach to CloudFront, ALB, or API Gateway .

GuardDuty vs. Security Hub vs. Inspector?

Answer: GuardDuty: Threat detection (logs/DNS). Security Hub: Central compliance dashboard. Inspector: Vulnerability scanning for EC2/Containers. Enable all for mature security .

How to secure cross-account access?

Answer: Use IAM Roles with `sts:AssumeRole` and External IDs. Federate identity via AWS Identity Center. Use SCPs to restrict dangerous actions .

Part 7: Integration & Serverless

SQS vs. SNS vs. EventBridge vs. Kinesis?

Answer: SQS: Pull-based queue. SNS: Pub/Sub fan-out. EventBridge: Event bus for AWS service events and SaaS. Kinesis: Real-time streaming with replay capability. Default to EventBridge for AWS events, SQS for buffering .

SQS Standard vs. FIFO?

Answer: Standard: At-least-once, best-effort ordering, unlimited throughput. FIFO: Exactly-once, strict ordering, lower throughput (3,000 msg/s). Use FIFO for financial transactions .

Lambda Concurrency: Reserved vs. Provisioned?

Answer: Reserved: Caps max executions to protect downstream. Provisioned: Keeps instances warm to eliminate cold starts. Use Provisioned for latency-sensitive APIs .

Step Functions vs. SQS-driven Lambda?

Answer: Step Functions: Visual workflow, error handling, retries, parallel branches. SQS+Lambda: Simpler, cheaper, fire-and-forget. Use Step Functions for complex orchestration (3+ steps) .

Part 8: Design Scenarios (Senior Level)**Design a highly available web application.**

Answer: CloudFront -> ALB (Multi-AZ) -> ASG (EC2/Fargate) -> Aurora Multi-AZ. Cache with ElastiCache. Static assets in S3. Use Route 53 Health Checks. Add WAF for security .

Design a serverless REST API for 100K RPS.

Answer: API Gateway (HTTP API) -> Lambda (Provisioned Concurrency) -> DynamoDB (On-Demand). Cache hot paths at API Gateway. Use X-Ray for tracing .

Design a Data Lake.

Answer: S3 (Raw/Curated zones) -> Glue Catalog -> Athena (Ad-hoc) -> Redshift (BI). Use Lake Formation for governance. Ingest via Kinesis/MSK .

Troubleshooting sudden slowness.

Answer: Check CloudWatch Dashboards (ALB latency, DB CPU). Use X-Ray traces to identify the bottleneck. Check for throttling in DynamoDB or Lambda concurrency limits .

Migrating a monolith to AWS.

Answer: Use the 6Rs. Often start with Rehost (Lift & Shift) using Application Migration Service, then refactor into containers (Strangler Fig pattern) over time. Use DMS for database migration .

"Tell me about a system you designed."

Answer: Use the STAR method (Situation, Task, Action, Result). Focus on the trade-offs you made (e.g., consistency vs. availability) and why you rejected other options. Mention Well-Architected pillars .

Preparation Tips for 2026

Focus on Trade-offs: Junior architects say "Use S3 because it's cheap." Senior architects say "Use S3 Standard for 30 days, then lifecycle to Glacier Deep Archive to save costs, accepting a 12-hour retrieval SLA" .

Whiteboard Practice: Be ready to draw a 3-tier app, serverless API, and multi-region DR setup from memory .

Latest Services: Be aware of new services like Amazon Q (AI assistant), Bedrock (GenAI), and updates to Control Tower.

Design an enterprise Agentic AI/RAG platform on AWS for a financial institution

This production-grade enterprise AI architecture isolates compute inside a private subnet, decoupling execution logic from foundation model inference while maintaining strict security and human-in-the-loop oversight.

Edge & Ingress Security

CloudFront, WAF & Shield: Intercepts traffic at the edge to block volumetric DDoS, SQL injection, and common web exploits before requests hit internal infrastructure.

Application Load Balancer (ALB): Terminates TLS, handles path-based routing, and forwards clean traffic directly to the private compute cluster.

Core Orchestration Layer (ECS/EKS Private Subnet)

AI Gateway: Manages rate limiting, request validation, API key rotation, and telemetry/logging for all incoming AI workloads.

Agent Orchestrator: Maintains conversation state, handles task decomposition, and drives step-by-step reasoning workflows using state machines or framework agents (e.g., LangChain/LlamaIndex/AutoGPT).

RAG Service: Manages text chunking, embedding generation, and vector retrieval pipelines for external context augmentation.

Execution & Data Integration

Guardrails: Scans incoming prompts and model outputs for PII leakage, toxicity, hallucination risks, and prompt injection attacks.

Knowledge Base (OpenSearch / Aurora): Provides hybrid search—combining OpenSearch Serverless vector embeddings with PostgreSQL (pgvector) structured relational context.

Tools Execution: Gives the agent agency via standardized protocol interfaces:

Risk & Trade APIs: External endpoints for domain-specific business execution.


SQL: Direct database queries for dynamic analytical retrieval.

MCP (Model Context Protocol): Standardized context integration across tools and context providers.

Inference & HITL Gateways

Amazon Bedrock (Claude / Llama / Titan): Managed, serverless Foundation Model inference inside AWS enterprise boundaries, eliminating data-retention concerns.

Human Approval (HITL): Enforces a "Human-in-the-Loop" circuit breaker for high-risk actions (e.g., executing trades or modifying databases) before final payload delivery to the **User**.



Here are 30+ targeted interview questions and answers focusing on AWS Generative AI, Amazon Bedrock, Amazon Q, Agentic AI, RAG architectures, and their associated monitoring and deployment resources as of 2026.

Part 1: Amazon Bedrock & Foundation Models (FM)

What is Amazon Bedrock and how does it differ from SageMaker?

Answer: Amazon Bedrock is a fully managed service that provides access to high-performing foundation models (FMs) from leading AI companies (Anthropic, Meta, Mistral, etc.) via a single API. It is designed for ease of use and rapid integration. SageMaker is a broader ML platform for building, training, and deploying custom ML models from scratch. Use Bedrock for off-the-shelf FMs; use SageMaker for custom model training or fine-tuning proprietary data .

What are the key features of Amazon Bedrock Agents?

Answer: Bedrock Agents allow you to build autonomous agents that can execute multi-step tasks across company systems and data sources. They automatically break down user requests into steps, call APIs, and manage conversation history. They integrate natively with Knowledge Bases for RAG .

Explain the concept of "Model Distillation" in Bedrock.

Answer: Model distillation allows you to create smaller, more cost-effective models by learning from larger, more complex teacher models. In Bedrock, this helps optimize performance and cost for specific enterprise use cases while maintaining high accuracy .

How does Amazon Bedrock handle data privacy?

Answer: By default, Amazon Bedrock does not use your input prompts or output generations to train the underlying foundation models. Data is encrypted at rest and in transit. You can also use VPC endpoints to keep traffic within your AWS network .

What is the difference between On-Demand and Provisioned Throughput in Bedrock?

Answer: On-Demand: Pay per token, no commitment, good for variable workloads. Provisioned Throughput: Purchase model units for a committed term (1-6 months) to guarantee consistent latency and throughput for production applications .

What are "Guardrails" in Amazon Bedrock?

Answer: Guardrails are configurable safety filters that help prevent generative AI applications from generating harmful or undesirable content. They can detect hate speech, sexual content, and prompt injection attacks, and can be applied across multiple models uniformly .

How do you evaluate FM performance in Bedrock?

Answer: Use Amazon Bedrock Model Evaluation, which provides automated metrics (accuracy, robustness, toxicity) and human-in-the-loop evaluation workflows. It helps compare models before deployment .

Part 2: RAG (Retrieval-Augmented Generation)

What is a Knowledge Base in Amazon Bedrock?

Answer: A fully managed RAG capability that connects FMs to your private data sources (S3, Confluence, Salesforce, etc.). It handles ingestion, chunking, vector embedding, and storage in an integrated vector store (like Aurora Serverless v2 or OpenSearch) without requiring you to build the pipeline manually .

What vector databases are supported natively by Bedrock Knowledge Bases?

Answer: Amazon Aurora PostgreSQL, Amazon OpenSearch Serverless, Pinecone, Redis Enterprise Cloud, and MongoDB Atlas. This flexibility allows you to choose the best engine for your scale and query requirements .

How does Bedrock handle "Chunking" in RAG?

Answer: Bedrock Knowledge Bases offer automatic chunking strategies (fixed-size, hierarchical, or semantic) during data ingestion. You can also define custom chunking logic using Lambda functions if the default strategies don't fit your document structure .

What is the role of Amazon Titan Embeddings in RAG?

Answer: Titan Embeddings models convert text into numerical vectors. These vectors are stored in the vector database. When a user queries, the query is embedded, and the system finds the most similar vectors (semantic search) to retrieve relevant context for the FM .

How do you improve retrieval accuracy in a RAG system?

Answer: Use Hybrid Search (combining keyword and semantic search), re-ranking models (to reorder retrieved documents by relevance), and optimize chunking strategies. Bedrock supports re-ranking via models like Cohere or Amazon Rerank .

Can Bedrock Knowledge Bases connect to real-time data?

Answer: Yes, via Lambda functions or API connectors within Agents. While Knowledge Bases primarily index static data, Agents can fetch real-time data from APIs during the execution phase .

Part 3: Agentic AI & Amazon Q**What is Amazon Q Developer?**

Answer: An AI-powered assistant for software development. It integrates with IDEs (VS Code, JetBrains) to provide code suggestions, generate unit tests, debug errors, and explain code. It can also help migrate legacy applications to newer frameworks .

How does Amazon Q Business differ from Q Developer?

Answer: Q Business is an enterprise-ready conversational assistant that connects to company data repositories (SharePoint, Slack, Jira) to answer business questions, summarize documents, and generate insights. Q Developer is focused on coding and IT operations .

What is "Agentic AI" in the context of AWS?

Answer: Agentic AI refers to systems that can perceive, reason, act, and learn autonomously to achieve goals. In AWS, this is powered by Bedrock Agents and Q Apps, which can plan multi-step workflows, call tools/APIs, and iterate based on feedback .

How do you build a custom agent in Bedrock?

Answer: Define an Action Group (API schema via OpenAPI/Swagger), select a Foundation Model, and optionally attach a Knowledge Base. The agent uses the FM to interpret user intent, map it to API calls, and return results .

What is the "CodeCatalyst" role in Agentic AI?

Answer: AWS CodeCatalyst (now largely integrated into Q Developer) accelerates development by automating repetitive tasks. In an agentic workflow, it can help generate infrastructure-as-code (IaC) templates or application scaffolding based on natural language prompts .

How does Amazon Q handle enterprise security?

Answer: Q respects existing IAM permissions and data access controls. It only retrieves and summarizes information the user is already authorized to see. It does not train on enterprise data unless explicitly configured for specific customization scenarios .

What are "Q Apps"?

Answer: No-code/low-code generative AI apps built on top of Amazon Q Business. Users can create custom assistants for specific teams (e.g., HR policy bot, Sales summary bot) without writing code, leveraging existing knowledge bases .

Part 4: Deployment, Monitoring, Logging & Evaluation

How do you deploy a Bedrock application to production?

Answer: Use AWS SAM or CDK to provision Bedrock resources (Agents, Knowledge Bases). Integrate with API Gateway for external access. Use CloudFront for global low-latency access. Ensure IAM roles follow least privilege for Bedrock actions .

How do you monitor Bedrock API usage and costs?

Answer: Use Amazon CloudWatch Metrics for token counts (input/output), latency, and error rates. Use AWS Cost Explorer with Bedrock-specific tags to track spend per model or application. Set up Budgets for alerts .

How do you log prompts and responses for auditing?

Answer: Enable Model Invocation Logging in Bedrock. Logs are sent to CloudWatch Logs or S3. You can capture input prompts, output generations, and metadata. Ensure sensitive data is masked before logging if required by compliance .

What is Amazon Bedrock Flows?

Answer: A visual interface to chain together multiple Bedrock components (Models, Knowledge Bases, Agents, Lambda functions) into a single workflow. It simplifies the deployment of complex GenAI pipelines without extensive coding .

How do you test for "Hallucinations" in a RAG app?

Answer: Use Bedrock Model Evaluation with "faithfulness" metrics. Implement a "grounding" check where the FM's answer is cross-referenced against the retrieved documents. If the answer contains info not in the context, flag it as a hallucination .

What is the role of AWS Lambda in GenAI architectures?

Answer: Lambda acts as the "glue" for preprocessing inputs, post-processing FM outputs, calling external APIs in Agents, and implementing custom business logic that FMs cannot perform directly (e.g., database writes) .

How do you secure GenAI applications against Prompt Injection?

Answer: Use Bedrock Guardrails to detect and block malicious prompts. Implement input validation in Lambda. Use separate contexts for system prompts and user inputs. Regularly test with adversarial prompts .

What is "Human-in-the-Loop" (HITL) in Bedrock?

Answer: A feature in Bedrock Evaluation and Agents where ambiguous or high-risk decisions are routed to human reviewers for validation before finalizing the output. This ensures quality and safety in critical workflows .

How do you manage versioning for Bedrock Agents?

Answer: Bedrock Agents support versioning. You can publish a draft agent as a version (v1, v2) and associate aliases (e.g., "prod", "dev") with specific versions. This allows safe rolling updates and A/B

testing .

What is the "Well-Architected Lens for GenAI"?

Answer: A specialized set of best practices for building GenAI workloads. It focuses on pillars like Responsible AI (fairness, transparency), Cost Optimization (token management), and Security (data privacy, guardrails). It extends the standard Well-Architected Framework [[30]].

How do you handle rate limiting in Bedrock?

Answer: Bedrock has default service quotas (TPS - Tokens Per Second). Use Application Auto Scaling for backend services. Implement retry logic with exponential backoff in your application code. Request quota increases via AWS Support if needed .

What is the role of Amazon EventBridge in Agentic AI?

Answer: EventBridge can trigger Agents or Bedrock Flows based on events (e.g., a new file in S3 triggers a summarization Agent). It enables event-driven GenAI workflows that react to real-time business changes .

Key Resources for 2026 Preparation:

AWS GenAI Chatbot Workshop: Hands-on labs for RAG and Agents.

Amazon Bedrock User Guide: For latest API changes and model support.

[AWS Well-Architected GenAI Lens: For architectural best practices.](#)

[Amazon Q Developer Documentation: For IDE integration and migration tools.](#)

Here are 30+ targeted interview questions and answers focusing on **Amazon EKS (Elastic Kubernetes Service)**, its integration with modern AWS services, security, networking, and its role in deploying AI/GenAI workloads in 2026.

Part 1: EKS Fundamentals & Architecture

What is Amazon EKS and how does it differ from ECS?

Answer: EKS is a managed Kubernetes service that allows you to run standard Kubernetes open-source tools. It is ideal for portability and complex microservices. ECS is AWS's proprietary container orchestration service, which is simpler to set up and deeply integrated with AWS APIs. Choose EKS for multi-cloud strategies or complex K8s ecosystems; choose ECS for simplicity and speed on AWS .

Explain the EKS Control Plane vs. Data Plane.

Answer: The Control Plane (API Server, etcd, Scheduler) is managed by AWS across multiple AZs for high availability. The Data Plane consists of worker nodes (EC2, Fargate, or Outposts) where your containers actually run. You are responsible for managing the data plane unless using Fargate .

What is the difference between EKS on EC2 and EKS on Fargate?

Answer: EKS on EC2: You manage node groups (patching, scaling, instance types). Offers full control, GPU support, and lower cost at scale. EKS on Fargate: Serverless. No node management. You pay per vCPU/memory second. Ideal for bursty workloads or when you want to eliminate ops overhead .

What is an EKS Add-on?

Answer: Managed software components that run on your cluster, such as VPC CNI, CoreDNS, and kube-

proxy. AWS manages their lifecycle, updates, and compatibility, reducing operational burden compared to self-managed Helm charts .

How does EKS handle version upgrades?

Answer: EKS supports rolling upgrades. You upgrade the Control Plane first, then update Node Groups. AWS provides a "Compatibility Matrix" to ensure your add-ons and applications support the new K8s version. Blue/Green deployments are recommended for critical production clusters .

What is EKS Auto Mode?

Answer: A newer feature (2024-2026) that automates node provisioning, scaling, and optimization. It automatically selects the best instance types, manages spot/on-demand mixing, and handles OS patching, reducing the need for manual Node Group configuration .

Part 2: Networking & Security in EKS

How does networking work in EKS (CNI)?

Answer: EKS uses the AWS VPC CNI plugin. Each Pod gets a native VPC IP address from the subnet. This allows Pods to communicate with other AWS services (like RDS or S3) using security groups and NACLs just like EC2 instances. It requires careful IP address planning .

What is the difference between Security Groups for Pods and Node Security Groups?

Answer: Node SGs: Protect the EC2 instance. Pod Security Groups: Allow you to define specific rules for individual Pods. This enables fine-grained network isolation between microservices running on the same node .

How do you secure access to the EKS API Server?

Answer: Use IAM Roles for Service Accounts (IRSA) for pod-level permissions. Restrict API server access via Kubernetes RBAC mapped to IAM identities. Use Private Access endpoint to keep traffic within the VPC. Enable CloudTrail for API auditing .

What is IRSA (IAM Roles for Service Accounts)?

Answer: IRSA allows Kubernetes Service Accounts to assume IAM Roles via OIDC federation. This provides least-privilege access to AWS resources (S3, DynamoDB) without embedding long-lived credentials in pods. It is the security best practice for EKS .

How do you expose an EKS application to the internet?

Answer: Use a Kubernetes Service of type LoadBalancer, which provisions an AWS ALB (Application Load Balancer) or NLB via the AWS Load Balancer Controller. For HTTP/HTTPS, ALB is preferred for path-based routing and WAF integration .

What is the AWS Load Balancer Controller?

Answer: A controller that runs in your EKS cluster and manages AWS ALBs and NLBs. It replaces the older "Ingress Controller" and allows you to define load balancer properties using Kubernetes Ingress or Service annotations .

Part 3: Storage & Stateful Workloads

How do you provide persistent storage to EKS Pods?

Answer: Use CSI (Container Storage Interface) Drivers. EBS CSI Driver for block storage (single pod access, high performance). EFS CSI Driver for shared file storage (multiple pods across nodes). FSx CSI Driver for high-performance computing (Lustre) or Windows workloads .

What is the difference between EBS and EFS in EKS?

Answer: EBS: Block storage, tied to a single AZ, higher IOPS, lower latency. Best for databases. EFS: Network file system, multi-AZ, shared by many pods. Best for content management or shared configs. EBS is generally cheaper for high-throughput single-instance workloads .

How do you handle secrets in EKS?

Answer: Avoid Kubernetes Secrets if possible. Use AWS Secrets Manager or Parameter Store mounted into pods via the Secrets Store CSI Driver. This allows automatic rotation and encryption at rest using KMS, keeping secrets out of etcd .

Part 4: EKS for AI/GenAI & ML Workloads

Why would you run GenAI workloads on EKS instead of SageMaker?

Answer: Use EKS when you need full control over the infrastructure, custom Kubernetes operators, or are running large-scale inference servers (like vLLM or TGI) that require specific GPU scheduling. Use SageMaker for end-to-end managed ML pipelines and easier model hosting .

How do you manage GPUs in EKS?

Answer: Install the NVIDIA Device Plugin and use GPU-enabled Instance Types (P5, G6, P4d). Use Karpenter or Cluster Autoscaler to provision GPU nodes on demand. Ensure you have the correct NVIDIA drivers installed via the EKS AMI .

What is Karpenter and why is it preferred over Cluster Autoscaler?

Answer: Karpenter is a next-gen node autoscaler for Kubernetes. It is faster, more flexible, and can provision nodes based on actual pod resource requests (including GPUs). It supports spot instances natively and reduces waste compared to the legacy Cluster Autoscaler .

How do you optimize costs for AI inference on EKS?

Answer: Use Spot Instances for fault-tolerant inference. Implement Model Serving frameworks (like Triton or vLLM) that support batching. Use Karpenter to right-size nodes. Scale down to zero using KEDA (Kubernetes Event-driven Autoscaling) when no requests are present .

Can EKS integrate with Amazon Bedrock?

Answer: Yes. Your EKS-hosted application can call Bedrock APIs using IRSA for authentication. You might use EKS to run a pre-processing pipeline or a RAG orchestrator that queries Bedrock Models and Knowledge Bases .

What is KServe or Seldon Core in EKS?

Answer: Open-source frameworks for serving ML models on Kubernetes. They handle canary deployments, A/B testing, and auto-scaling for ML models. They are often used in EKS to manage complex inference lifecycles that Bedrock doesn't cover (e.g., custom proprietary models) .

Part 5: Monitoring, Logging & Observability

How do you monitor EKS clusters?

Answer: Use Amazon CloudWatch Container Insights for metrics (CPU, Memory, Pod status). Use AWS Distro for OpenTelemetry (ADOT) to collect traces and logs. Integrate with Prometheus/Grafana (via AMP - Amazon Managed Prometheus) for detailed K8s-native monitoring .

How do you centralize logs from EKS?

Answer: Use Fluent Bit (daemonset) to ship logs from all nodes/pods to CloudWatch Logs or OpenSearch. For high-volume logs, consider streaming to S3 via Kinesis Data Firehose for cost-effective long-term retention .

What is X-Ray and how does it help in EKS?

Answer: AWS X-Ray provides distributed tracing. In EKS, it helps visualize the flow of requests across microservices, identifying latency bottlenecks in complex GenAI chains (e.g., API Gateway -> Lambda -> EKS Pod -> Bedrock) .

How do you debug a "CrashLoopBackOff" in EKS?

Answer: Check `kubectl describe pod` for events. Check `kubectl logs` for application errors. Common causes: missing config maps/secrets, insufficient memory (OOMKilled), or application startup failures. Use CloudWatch Logs for historical context .

Part 6: Advanced EKS & GitOps

What is GitOps and how is it implemented in EKS?

Answer: GitOps uses Git as the single source of truth for infrastructure and application deployment. Tools like ArgoCD or Flux are installed in EKS to sync the cluster state with Git repositories. This ensures auditability and rollback capabilities .

What is EKS Anywhere?

Answer: Allows you to run EKS clusters on-premises or in other clouds using the same AWS management plane. Useful for hybrid cloud strategies where data residency or latency requires local processing .

How do you handle multi-tenancy in EKS?

Answer: Use Namespaces for logical separation. Apply Resource Quotas and Limit Ranges to prevent noisy neighbors. Use Network Policies to isolate traffic between tenants. For strong isolation, use separate clusters or virtual clusters (vClusters) .

What is the role of AWS App Mesh in EKS?

Answer: App Mesh is a service mesh that provides application-level networking. It enables observability, traffic control, and security (mTLS) between microservices without changing code. It integrates with Envoy proxies injected into pods .

How do you secure the supply chain in EKS?

Answer: Use ECR (Elastic Container Registry) with image scanning. Sign images using Cosign. Use Admission Controllers (like OPA Gatekeeper or Kyverno) to enforce policies (e.g., "only signed images from our ECR repo can run"). Enable Image Vulnerability Scanning in ECR .

Key Resources for 2026 EKS Preparation:

EKS Best Practices Guide: Official AWS documentation for security, networking, and operations.

AWS Containers Blog: For latest updates on Karpenter, EKS Auto Mode, and GenAI on K8s.

Well-Architected Lens for Kubernetes: Specific guidance on reliability and cost for EKS.

Amazon EKS Workshop: Hands-on labs for IRSA, Fargate, and GitOps.

The Senior Solutions Architect - Developer Transformation role at Amazon is highly specialized. It sits at the intersection of Enterprise Modernization, Developer Experience (DevEx), and Strategic Consulting.

You are not just building systems; you are convincing CTOs and VP-level engineering leaders to change how their teams build software. You need to demonstrate expertise in Legacy Modernization, Mainframe Offloading, CI/CD at Scale, and Cultural Change Management.

Here are 30+ targeted questions and answers specifically for this role, focusing on what was missing in the previous lists.

Part 1: Legacy Modernization & Mainframe Offloading (Critical for this Role)

What is the "Strangler Fig Pattern" and when do you recommend it? [monolithic to microservices]

Answer: A pattern where you gradually replace specific pieces of functionality from a monolith with new microservices. Over time, the new services "strangle" the old system until it can be decommissioned. I recommend it for large, critical monoliths where a "big bang" rewrite is too risky. It allows for incremental value delivery and risk mitigation .

Case Scenario

Legacy System: A 10-year-old Java/Spring Boot monolith running on Amazon EC2, backed by a single Amazon RDS for Oracle database. It handles User Auth, Product Catalog, and Order Processing.

Goal: Extract the Order Processing domain into a modern, scalable, cloud-native microservice without disrupting the business ("Big Bang" rewrite is forbidden).

Target Architecture: Serverless/Microservice using Amazon API Gateway, AWS Lambda, and Amazon DynamoDB.

Phase 1: Establish the Facade (The "Vine")

Before writing any new code, we must intercept all traffic to control the routing.

Deploy Amazon API Gateway: Place an API Gateway (REST or HTTP API) in front of the existing monolith.

Configure Catch-All Routing: Initially, configure a proxy resource (`/api/{proxy+}`) in API Gateway that integrates via HTTP to the existing Application Load Balancer (ALB) in front of the EC2 monolith.

Result: The external world now talks to API Gateway, but the monolith still does 100% of the work. This is our baseline.

Phase 2: Build the New Microservice (The "New Growth")

We build the new Order Service in parallel, treating it as a greenfield project with modern DevEx practices.

Compute: Create an AWS Lambda function (e.g., in Node.js or Python) to handle `GET /orders` and `POST /orders`.

Data: Provision an Amazon DynamoDB table for the new service. This establishes a clear **Bounded Context**; the new service owns its data.

CI/CD: Set up an AWS CodePipeline with AWS CodeBuild and AWS CodeDeploy. This ensures the new service can be deployed independently and safely, a key metric for Developer Transformation.

Anti-Corruption Layer (ACL): The Lambda function includes logic to translate the new, clean API contracts into whatever format the legacy system might still need during the transition.

Phase 3: Data Migration & Synchronization (The Critical Step)

You cannot have two systems of record. We must migrate historical data and handle the transition safely.

Historical Migration: Use AWS Database Migration Service (DMS) to perform a one-time full load of historical order data from RDS Oracle to DynamoDB (or an intermediate Aurora PostgreSQL if relational integrity is strictly required).

Dual-Write Strategy (Transition Period):

For a brief period, the API Gateway routes `POST /orders` to both the new Lambda function and the legacy monolith (using API Gateway's integration request mapping or an EventBridge rule).

Better Alternative:* Route the `POST` **only to the new Lambda. The Lambda writes to DynamoDB, then publishes an `OrderCreated` event to Amazon EventBridge. A lightweight adapter in the monolith listens to this event and updates the legacy Oracle DB. This ensures the monolith remains functional for legacy reporting while the new service becomes the system of record.

Phase 4: Traffic Shifting (The "Strangling")

Now we gradually shift traffic from the monolith to the new service using API Gateway routing rules.

Update API Gateway Routes:

Create a specific route for `/api/orders/{orderId}` and `/api/orders` (POST).

Change the integration for these specific routes from the HTTP ALB (monolith) to the new Lambda Function.

Leave all other routes (e.g., `/api/products`, `/api/users`) pointing to the monolith.

Canary Deployment: Use AWS CodeDeploy with Lambda Aliases to shift traffic gradually (e.g., 10% to the new service, 90% to the monolith).

Observability: Enable AWS X-Ray across API Gateway, the new Lambda, and the legacy EC2 app (via the X-Ray daemon/SDK). Use Amazon CloudWatch Dashboards to compare error rates and latency between the old and new paths.

This is 100% on the new service, and the legacy Oracle database no longer has active Order writes.

Code Removal: Developers remove the Order Processing code, controllers, and dependencies from the Java monolith codebase. This reduces the monolith's complexity and build time (a huge win for Developer Experience).

Database Cleanup: Archive the old Order tables in RDS Oracle to Amazon S3 (via DMS or native export) for compliance, then drop the tables to free up resources and licensing costs.

Facade Cleanup: The API Gateway is now the permanent API for the system. The ALB in front of the monolith now only serves the remaining domains (e.g., Product Catalog), which will be the

next candidates for the Strangler pattern.

How do you approach Mainframe Modernization on AWS?

Answer: I use the AWS Mainframe Modernization Service or partner solutions (like Blu Insights). The strategy depends on the goal: Rehost (lift-and-shift using Micro Focus or Raincode), Refactor (convert COBOL to Java/Python using automated tools), or Rebuild (rewrite using cloud-native services like Lambda/DynamoDB). The key is to decouple data first using CDC (Change Data Capture) via DMS .

What is the role of AWS Application Migration Service (MGN)?

Answer: MGN is used for "Lift and Shift" (Rehosting). It replicates servers at the block level to AWS with minimal downtime. For a Developer Transformation role, I position MGN as a *first step* to get out of the data center quickly, followed by a *second phase* of refactoring/modernization once on AWS .

How do you handle database modernization from Oracle/SQL Server to Aurora?

Answer: Use AWS Schema Conversion Tool (SCT) to convert schema and code (stored procedures). Use Database Migration Service (DMS) for continuous data replication. I often recommend a "dual-run" period where both databases are updated to ensure data integrity before cutover. Moving to Aurora Serverless v2 can also reduce operational overhead .

What are the biggest cultural challenges in legacy modernization?

Answer: Resistance to change, lack of cloud skills, and fear of job loss. As a SA, I address this by establishing Centers of Excellence (CoE), providing hands-on training, and demonstrating quick wins. I emphasize that modernization frees developers from mundane ops tasks to focus on innovation .

Part 2: Developer Experience (DevEx) & Internal Developer Platforms (IDP)

What is an Internal Developer Platform (IDP) and why do enterprises need it?

Answer: An IDP is a self-service layer that abstracts infrastructure complexity. It allows developers to provision environments, deploy code, and monitor apps without waiting for Ops. On AWS, this is built using Service Catalog, Backstage, or AWS Proton. It reduces cognitive load and accelerates time-to-market .

How does AWS Proton fit into Developer Transformation?

Answer: AWS Proton allows platform engineers to define standardized templates for infrastructure (ECS/EKS/Lambda) and CI/CD pipelines. Developers simply select a template and deploy their code. It enforces best practices and security guardrails automatically, ensuring consistency across hundreds of teams .

What metrics do you use to measure Developer Productivity?

Answer: I use the DORA Metrics (DevOps Research and Assessment):

Deployment Frequency: How often do you deploy?

Lead Time for Changes: Time from commit to production.

Change Failure Rate: Percentage of deployments causing failure.

Time to Restore Service: How long to recover from failure. These metrics help quantify the impact of transformation efforts .

How do you implement "Golden Paths" or "Paved Roads"?

Answer: Create pre-approved, secure, and documented templates for common architectures (e.g., "Serverless API," "Microservice on EKS"). Make these paths the easiest way to build. If developers go off-road, they bear the burden of security and compliance reviews. This guides behavior without forcing it .

What is the role of Backstage in AWS environments?

Answer: Backstage is an open-source developer portal. On AWS, it integrates with Service Catalog and CodePipeline to provide a unified UI for developers to discover APIs, create new services, and view documentation. It centralizes the developer experience .

Part 3: CI/CD at Enterprise Scale

Compare AWS CodePipeline vs. Jenkins/GitLab CI.

Answer: CodePipeline is fully managed, integrates natively with AWS services (CodeBuild, CodeDeploy), and requires no server maintenance. Jenkins/GitLab offer more flexibility and plugin ecosystems but require significant operational overhead. For enterprise transformation, I often recommend migrating to CodePipeline or GitHub Actions to reduce ops burden and improve security .

How do you implement Blue/Green Deployments on AWS?

Answer:

EC2/ECS: Use CodeDeploy with an ALB to shift traffic gradually.

Lambda: Use CodeDeploy with weighted aliases.

EKS: Use Argo Rollouts or Flagger with Istio/App Mesh.

Route 53: Use weighted routing policies for DNS-level switching. This ensures zero-downtime deployments and easy rollback .

What is GitOps and how do you enforce it in a large organization?

Answer: GitOps uses Git as the single source of truth for infrastructure and apps. I enforce it by using ArgoCD or Flux in EKS. Policies are defined in Git (e.g., Terraform/CloudFormation).

Any change must go through a Pull Request with peer review and automated checks before being applied to the cluster. This provides auditability and prevents drift .

How do you secure the CI/CD pipeline?

Answer: Implement SBOM (Software Bill of Materials) generation. Scan images in ECR for vulnerabilities. Use IAM Roles for Service Accounts (IRSA) so pipelines have least-privilege access. Sign artifacts using Sigstore/Cosign. Integrate Security Scanning (Checkov, tfsec) into the build stage .

What is the "Shift Left" security model?

Answer: Integrating security checks early in the development lifecycle (IDE, Commit, Build) rather than at the end. Tools like Amazon CodeGuru Security or Static Analysis in CodeBuild catch vulnerabilities before they reach production. This reduces the cost and effort of fixing security issues .

Part 4: Strategic Consulting & Leadership Principles

How do you handle a customer who says "Cloud is too expensive"?

Answer: I conduct a Well-Architected Review focused on Cost Optimization. I look for idle resources, over-provisioned instances, and lack of reserved capacity. I introduce FinOps practices: tagging strategies, budget alerts, and showback/chargeback models. Often, the issue is not cloud cost but *waste* due to poor architecture .

Describe a time you influenced a technical decision against a customer's initial preference.

Answer: (Use STAR Method). Example: Customer wanted to lift-and-shift a monolithic Oracle DB to EC2. I demonstrated that moving to Aurora PostgreSQL would reduce licensing costs by 80% and improve availability. I built a PoC to prove performance parity. Result: Customer adopted Aurora and saved \$2M/year .

How do you assess an organization's "Cloud Maturity"?

Answer: I use the AWS Cloud Adoption Framework (CAF). I evaluate six perspectives: Business, People, Governance, Platform, Security, and Operations. I identify gaps (e.g., "Strong Platform but weak Governance") and create a roadmap to address them. This holistic view ensures sustainable transformation .

What is the role of a "Center of Excellence" (CoE)?

Answer: A CoE is a cross-functional team that defines best practices, creates reusable assets (templates, pipelines), and trains other teams. It acts as the internal consultancy. I help customers establish CoEs to scale cloud knowledge beyond the initial pilot team .

How do you align technical transformation with business goals?

Answer: I start by understanding the business KPIs (e.g., "Reduce time-to-market by 50%"). Then, I map technical initiatives (e.g., "Implement CI/CD," "Adopt Microservices") directly to those KPIs. I avoid tech-for-tech's sake. Every architectural recommendation must have a clear business value proposition .

Part 5: Advanced Modernization Patterns

What is Event-Driven Architecture (EDA) and why is it key to modernization?

Answer: EDA decouples services using events (SNS/SQS/EventBridge). It allows legacy systems to emit events without knowing who consumes them. This enables gradual decomposition of monoliths. New services can subscribe to legacy events, allowing them to coexist and eventually replace legacy logic .

How do you handle distributed transactions in a microservices architecture?

Answer: Avoid two-phase commits. Use the Saga Pattern (Choreography or Orchestration). In AWS, Step Functions is ideal for orchestrating sagas, managing retries, and compensating transactions if a step fails.

This ensures data consistency across services without tight coupling

What is the "Anti-Corruption Layer" (ACL)?

Answer: An ACL is a translation layer between a new system and a legacy system. It prevents legacy design flaws and data models from "corrupting" the new clean architecture. It translates requests/responses between the two worlds. This is critical when integrating modern microservices with mainframes .

How do you modernize .NET Framework applications?

Answer: Use AWS App2Container to containerize .NET apps. Run them on EKS or ECS. For newer versions, migrate to .NET Core/.NET 6+ and run on Lambda or Fargate. Use Windows Containers on ECS/EKS if full framework compatibility is needed .

What is the role of Amazon MQ in modernization?

Answer: Amazon MQ (ActiveMQ/RabbitMQ) is used when applications require standard protocols (JMS, AMQP) that SNS/SQS don't support. It's often a stepping stone: lift-and-shift existing JMS apps to Amazon MQ, then gradually refactor to native AWS services (SQS/SNS) later .

Part 6: AI/GenAI in Developer Transformation

How does Amazon Q Developer accelerate transformation?

Answer: Q Developer can analyze legacy code (Java, COBOL, .NET) and suggest modernizations. It can generate unit tests, document code, and even convert code between languages. I recommend it to speed up the "Refactor" phase of the 6Rs .

How do you govern GenAI usage in an enterprise?

Answer: Establish an AI Center of Excellence. Use Amazon Bedrock Guardrails to enforce safety. Implement Identity Federation so only authorized users can access models. Monitor usage with CloudTrail and Cost Allocation Tags. Create policies for data privacy (no PII in prompts) .

What is "Agentic CI/CD"?

Answer: Using AI Agents to manage pipelines. For example, an agent detects a failed build, analyzes logs, suggests a fix, and creates a PR. Or an agent auto-scales infrastructure based on predictive traffic. This moves DevOps from "Automated" to "Autonomous" .

How do you measure the ROI of GenAI in development?

Answer: Track metrics like Code Acceptance Rate (how much AI-generated code is merged), Reduction in Boilerplate Coding Time, and Defect Density. Use Amazon Q Business analytics to see adoption rates across the organization .

What is the biggest risk of adopting GenAI in coding?

Answer: Intellectual Property (IP) leakage and Security Vulnerabilities in generated code. Mitigate by using enterprise-grade tools like Amazon Q Developer (which doesn't train on your code), implementing strict code review processes, and using static analysis tools to scan AI-generated code .

Key Preparation Tips for This Specific Role:

Know the "6 Rs" of Migration Cold: Rehost, Replatform, Refactor, Repurchase, Retire, Retain. Be ready to give examples of each.

Focus on "People" and "Process": This role is 50% technology and 50% change management. Talk about Training, CoEs, and Culture.

Master the "Why": Don't just say *what* service to use. Explain *why* it solves a business problem (e.g., "We use Proton to reduce developer onboarding time from 2 weeks to 2 hours").

Read the "[AWS Cloud Adoption Framework](#)" (CAF): This is the bible for this role. Understand the 6 perspectives deeply.

Prepare STAR Stories: Have 3-5 strong stories about leading complex transformations, dealing with resistant stakeholders, and delivering measurable business value.

Explain about aws migration services and how can it help to on premises or other cloud services to aws

Moving from on-premises or other clouds (Azure, GCP) to AWS is not a "one-size-fits-all" process. AWS provides a specialized portfolio of tools categorized by the 6 Rs of Migration strategy: *Rehost, Replatform, Refactor, Repurchase, Retire, and Retain*.

Here is a breakdown of the key AWS Migration Services and their specific use cases.

Discovery & Planning: "What do we have?"

Before moving anything, you need to know what exists.

AWS Migration Hub: A central dashboard that aggregates data from discovery tools. It provides a single view of your migration progress across all accounts and regions.

* Application Discovery Service (ADS): Agentless or agent-based discovery that maps dependencies between servers. *Crucial for identifying which monolith components talk to which databases.*

* Migration Evaluator (formerly TCO Calculator): Uses historical usage data to recommend right-sized AWS instances and estimate cost savings compared to on-prem.

Data Transfer: "How do we get the bits there?"

AWS DataSync: Online transfer service. Best for moving active data (files, NFS/SMB shares) over the network. It handles incremental syncs and verification.

AWS Snow Family (Snowcone, Snowball, Snowmobile): Offline physical devices. Used when network bandwidth is insufficient or too expensive for Petabytes of data. You ship the device to AWS; they plug it into S3.

Amazon S3 Transfer Acceleration: Uses CloudFront edge locations to speed up long-distance uploads to S3.

Server & VM Migration (Rehost / Lift-and-Shift)

AWS Application Migration Service (MGN): (The Flagship Tool)

How it works: It installs a lightweight agent on source servers (VMware, Hyper-V, Azure, GCP, Physical). It performs continuous block-level replication to AWS.

- * Key Benefit: Minimal Downtime. You can launch test instances in AWS anytime. When ready for cutover, you stop the source, replicate the final delta, and launch the target instance in minutes.

- * Use Case: Moving large fleets of EC2-compatible VMs without refactoring code.

Database Migration (Replatform / Refactor)

AWS Database Migration Service (DMS):

How it works: Creates a replication instance that connects to source and target databases.

It supports Continuous Data Replication (CDC).

Key Benefit: Heterogeneous Migrations. It can migrate from Oracle/SQL Server to Amazon Aurora (PostgreSQL/MySQL compatible). It keeps the source DB running during migration, allowing for near-zero downtime cutover.

Schema Conversion: Often paired with AWS Schema Conversion Tool (SCT), which converts database schema objects (stored procedures, triggers) from one engine to another.

Amazon RDS Blue/Green Deployments: For minor version upgrades or parameter changes within RDS, this creates a safe staging environment to test before switching traffic.

Mainframe & Legacy Modernization (Refactor / Replatform)

AWS Mainframe Modernization Service:

How it works: Provides a managed runtime environment for mainframe applications (COBOL, PL/I) using partners like Blu Insights or Micro Focus.

Use Case: Customers who want to exit the data center but cannot rewrite 30-year-old COBOL code immediately. It allows "Rehosting" mainframes on AWS infrastructure.

AWS Partner Solutions (e.g., Raincode, Astera): Often integrated via the Migration Hub to automate code conversion from COBOL/PL-I to Java/.NET.

Application & Code Modernization (Refactor)

AWS App2Container (A2C):

How it works: Analyzes running Java or .NET applications on-prem, identifies dependencies, and packages them into Docker containers. It generates Kubernetes manifests (for EKS) or ECS task definitions.

Use Case: Moving monolithic Java/.NET apps to containers without manual Dockerfile creation.

Amazon Q Developer (formerly CodeWhisperer + more):

How it works: AI-powered assistant that can analyze legacy code, suggest refactors, generate unit tests, and even help translate code from one language to another (e.g., Java to Python).

Use Case: Accelerating the "Refactor" phase of the 6 Rs by reducing manual coding effort.

End-of-Life & Decommissioning (Retire)

AWS Cost Explorer & Trusted Advisor: Identify unused resources post-migration to ensure you aren't paying for both on-prem and AWS simultaneously longer than necessary.

Strategic Framework: The 6 Rs in Practice

| Strategy | AWS Service(s) | Use Case |

| :--- | :--- | :--- |

| Rehost (Lift & Shift) | Application Migration Service (MGN) | Quick move to AWS to exit data center lease. No code changes. |

| Replatform (Lift & Reshape) | RDS, Elastic Beanstalk, ECS | Move DB to Aurora, app to managed containers. Minor optimizations. |

| Refactor (Modernize) | App2Container, Lambda, EKS, Amazon Q | Break monolith into microservices. Change code/architecture. |

| Repurchase (Drop & Shop) | Salesforce, Workday, SaaS | Move from custom CRM to Salesforce. Migrate data via API/DataSync. |

| Retire | Cost Explorer | Turn off unused servers/apps discovered during assessment. |

| Retain | Snowball, Direct Connect | Keep sensitive/legacy apps on-prem for now, but connect to AWS. |

Interview Answer Example: "How would you migrate a critical Oracle DB and Java Monolith from On-Prem to AWS with minimal downtime?"

"I would approach this using a hybrid strategy focusing on minimal downtime and risk mitigation:

>

1. Discovery: First, I'd use Application Discovery Service to map dependencies between the Java app and the Oracle DB.

2. Database Migration: I'd use AWS SCT to convert the Oracle schema to Amazon Aurora PostgreSQL. Then, I'd use AWS DMS to perform a full load followed by Continuous Data Replication (CDC). This keeps the on-prem Oracle and Aurora in sync, allowing us to cut over at any moment with seconds of downtime.

3. Application Migration: For the Java monolith, if we are doing a 'Lift and Shift' first, I'd use AWS Application Migration Service (MGN). It replicates the VM block-by-block. We can test the app in AWS repeatedly without impacting production.

4. Cutover: Once the app team validates the Aurora connection strings and functionality in the MGN test instance, we schedule a maintenance window. We stop writes to Oracle, let DMS flush the final CDC changes, update DNS (Route 53) to point to the new AWS ALB, and start the app.

5. Future State: Post-migration, I'd recommend using App2Container or Amazon Q Developer to gradually refactor the monolith into microservices on EKS, leveraging the Strangler Fig pattern we discussed earlier."