

Stanford LLM Lecture Series (watch in order)

- Transformer Fundamentals — <https://www.youtube.com/watch?v=Ub3GoFaUcds>
- Transformer Models + Practical Tricks — <https://www.youtube.com/watch?v=yT84Y5zCnaA>
- Transformers → Large Language Models — <https://www.youtube.com/watch?v=Q5baLehv5So>
- How LLMs Are Trained — https://www.youtube.com/watch?v=VIA_jt_3Qc4
- Tuning & Adaptation (fine-tuning, etc.) — https://www.youtube.com/watch?v=PmW_TMq3l0I
- Reasoning in LLMs — <https://www.youtube.com/watch?v=k5Fh-UgTuCo>
- Agentic LLMs (tools, planning, workflows) — <https://www.youtube.com/watch?v=h-7S6HNq0Vg>
- Evaluation: what "good" really means — <https://www.youtube.com/watch?v=8fNP4N46RRo>
- Recap + What's Trending Now — <https://www.youtube.com/watch?v=Q86qzJ1K1Ss>

Pair the theory with hands-on labs

To make this path more practical and well-rounded, I'd pair the Stanford theory with these free resources:

- Microsoft — Generative AI for Beginners: <https://github.com/microsoft/generative-ai-for-beginners>
- NVIDIA Developer Program (tools, SDKs, learning): <https://developer.nvidia.com/join-nvidia-developer-program>
- Google Cloud — ML & AI Training tracks: <https://cloud.google.com/learn/training/machinelearning-ai>