Microsoft Azure

A leader's guide to building high-impact generative AI apps

Insights from AI pioneers for setting your organisation up to build better, more scalable generative AI applications

Executive summary

Generative AI is transforming how businesses operate, innovate and serve customers – and it's doing so at unprecedented speed. At the same time, many leaders are still working out how to advance from experimentation to enterprise-grade solutions that drive real business value.

In this e-book, we'll share insights from AI-pioneering organisations and customers both inside and outside Microsoft on how to set up your generative AI projects for success. This guide will help you navigate the complexities of building high-impact generative AI apps – providing a roadmap from ideation to deployment that focuses on the full flexibility and power of custom-built AI applications.

You'll discover how to:

- → Align AI projects with business value
- → Choose use cases that address key business needs
- → Assemble a cross-functional team
- → Prepare an AI-ready foundation
- → Help your organisation embrace a test and learn culture

Who is this e-book for?

This e-book is for technical and business leaders who are looking to infuse AI into their new and existing applications to enhance customer experiences and drive business value.

In addition to Azure solutions, Microsoft offers a range of solutions to help build customisable AI applications. Whether through Visual Studio Code, GitHub Copilot or other AI-powered tools, there are flexible options available to meet your needs and drive AI innovation across your organisation.

Estimated reading time: **20 minutes**

Contents

Introduction

We're living through one of the most significant platform shifts in decades – the transition to an AI-powered future. Just as cloud computing and mobile technologies reshaped how we work and innovate, generative AI is now redefining what's possible for organisations of every size and industry.

The **AI platform shift** is accelerating a new wave of application development. Organisations are now building apps that use generative AI to engage users in more natural ways, adapt dynamically to context and deliver new kinds of value – faster. Whether enhancing existing applications or creating entirely new solutions grounded in their own data, businesses are reimagining how their software works and what it can achieve.

This isn't just a technology transformation – it's a business one. Leaders are rethinking how they define success, how teams work together and how to responsibly scale solutions. But many still face uncertainty about how to get started or move from pilots to production.

That's where this e-book comes in.

Drawing from in-depth research and conversations with AI-pioneers – including those building generative AI apps on Microsoft Azure – this guide offers practical insights for leaders looking to build high-impact, custom generative AI apps. You'll find real-world examples, key decision points and guidance on everything from identifying use cases and preparing your data to standing up cross-functional teams and establishing an iterative, responsible development process.

Whether you're experimenting with your first solution or scaling AI across your enterprise, this resource is designed to help you take the next step – with clarity, purpose and confidence.

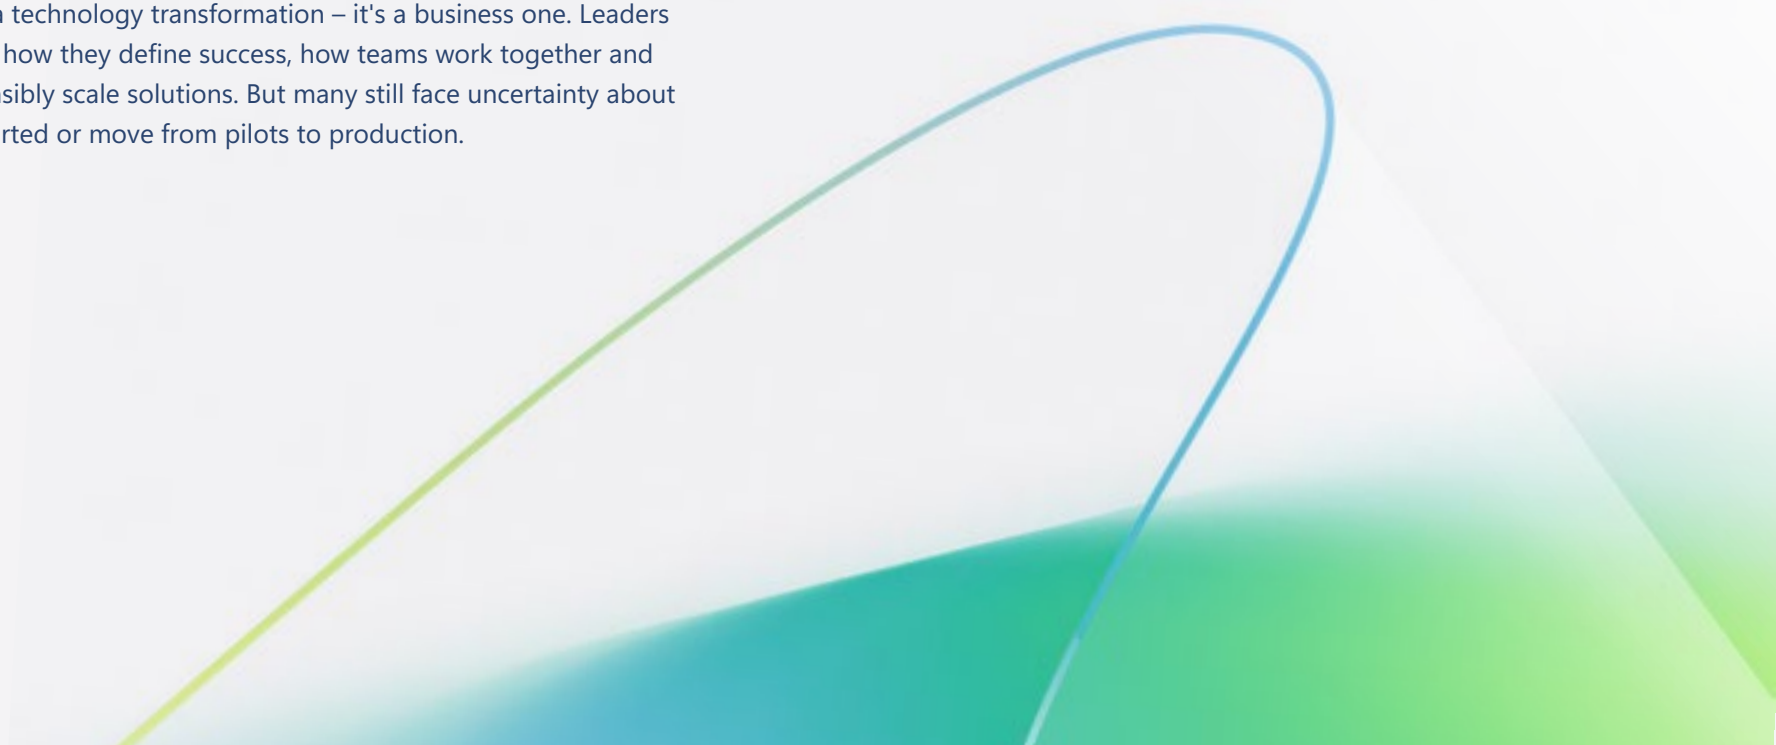

What leaders consider when building generative AI apps

Insight No. 1

Leading AI adopters redefine success beyond ROI

AI pioneers find that traditional ROI metrics don't capture AI's full value.

A common theme among leaders is that traditional return on investment (ROI) metrics capture only part of AI's true value. For example, early investments in Internet technology might have yielded minimal ROI at first, but they enabled organisations to increase efficiency, compete more effectively in the market, deliver product, service and business model innovation, and better understand and serve their customers in the longer term.

AI is similar. Like the Internet, AI comprises a range of technologies used for a range of use cases to meet a range of business objectives across a range of industries. As a result, leaders tend to quantify the impact of AI value beyond ROI, in terms of specific use cases and in a way that resonates with the C-suite and their boards of directors.

Which is the most important business outcome that your organisation is trying to achieve from AI intiatives?

(Percentage of respondents)

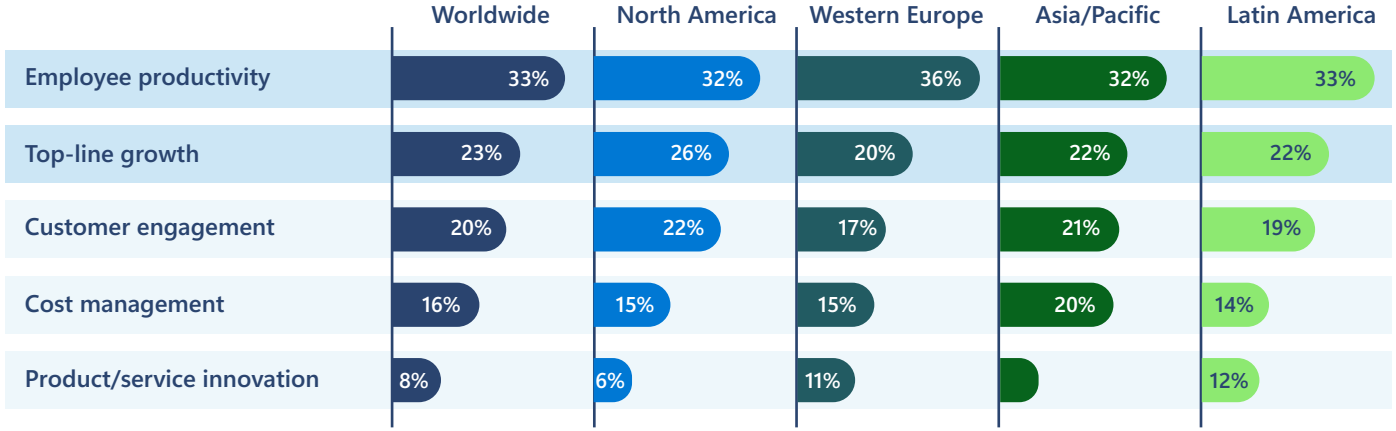

Source: IDC Business Opportunity of AI 2024

How AI-pioneering organisations approach ROI differently

Leading AI adopters take a nuanced approach to measuring success, expanding beyond financial metrics to assess AI's broader impact. These organisations focus instead on industry-specific outcomes, stakeholder impact and cost savings in addition to revenue growth.

Common choices for measures of success we heard from these pioneering AI leaders include:

- **Output quality:** How accurate and useful are the results from the custom-built generative AI solution and is it better than the alternative (e.g. better than a manual process or a pre-existing solution)?
- **User adoption:** How many users are using the solution?
- **User engagement:** How frequently are the intended users using the solution?
- **User satisfaction:** How happy are the users with the solution?
- **Time to completion:** How long does it take users to perform a task with the solution and how does that compare to how long it used to take with a previous solution or manual process?
- **Breadth of impact:** How many parts of the organisation does this solution serve, just one business unit or many?

According to IDC organisations realise a return on their AI investments within 14 months of deployment.[3]

Key takeaways from AI leaders:

- **Expanding AI value categories:** AI delivers measurable benefits in customer service, risk management and operational efficiency that go beyond immediate revenue impact.
- **Recognising AI's ROI timeline:** While financial ROI may take 12 to 18 months or more to materialise, early benefits – such as increased efficiency and improved decision-making – emerge much sooner.
- **Measuring AI beyond financials:** Organisations track success through innovation, employee skill-building and brand perception.

Traditional financial metrics such as ROI remain critical nonetheless. An AI investment yielding £2 million in savings or net-new revenue from a £1 million expenditure demonstrates clear value.

"Every investment must be justified. If AI can replace an existing cost, its value is immediately clear."

Vice President
Business & Technology Transformation, Financial Services, UK[2]

"Many companies struggle to define AI's ROI. We developed our own framework, which considers portability across business units, innovation potential and investor perception – beyond just financials."

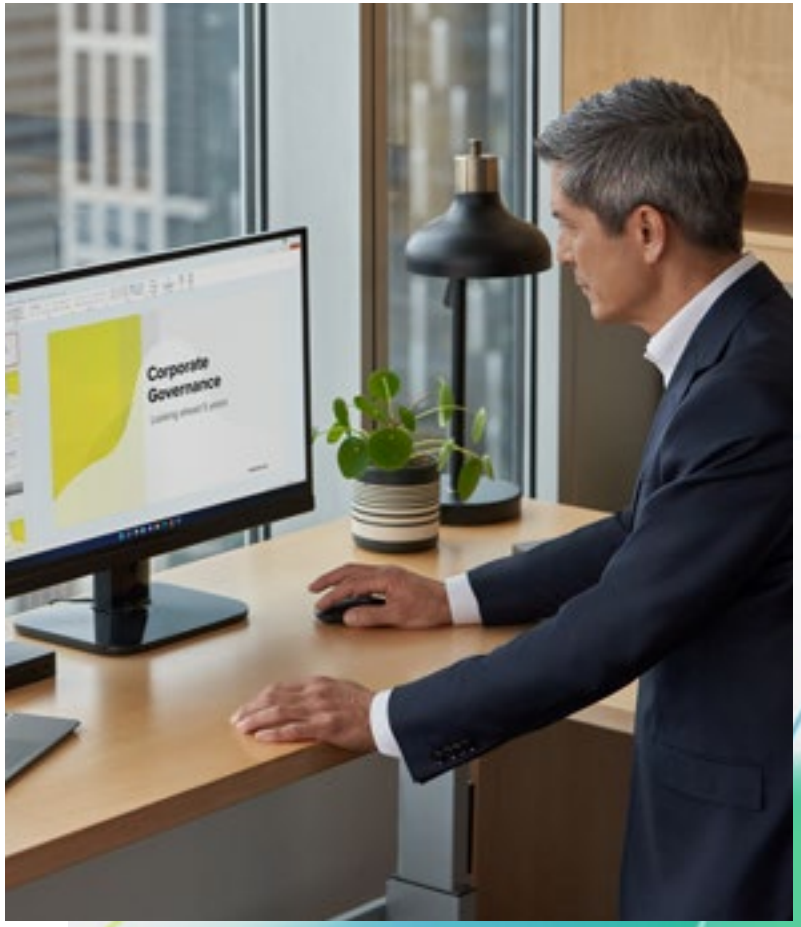

Chief AI Officer
Insurance, UK[2]

AI leaders recognise that **the value of AI extends beyond traditional ROI calculations**. While financial returns are important, they also consider factors like growth potential, risk management and operational efficiency. These elements help determine **how AI initiatives contribute to long-term success and competitive advantage**.

By evaluating AI through a variety of lenses, leaders ensure that their AI investments align with broader business goals and deliver meaningful results.

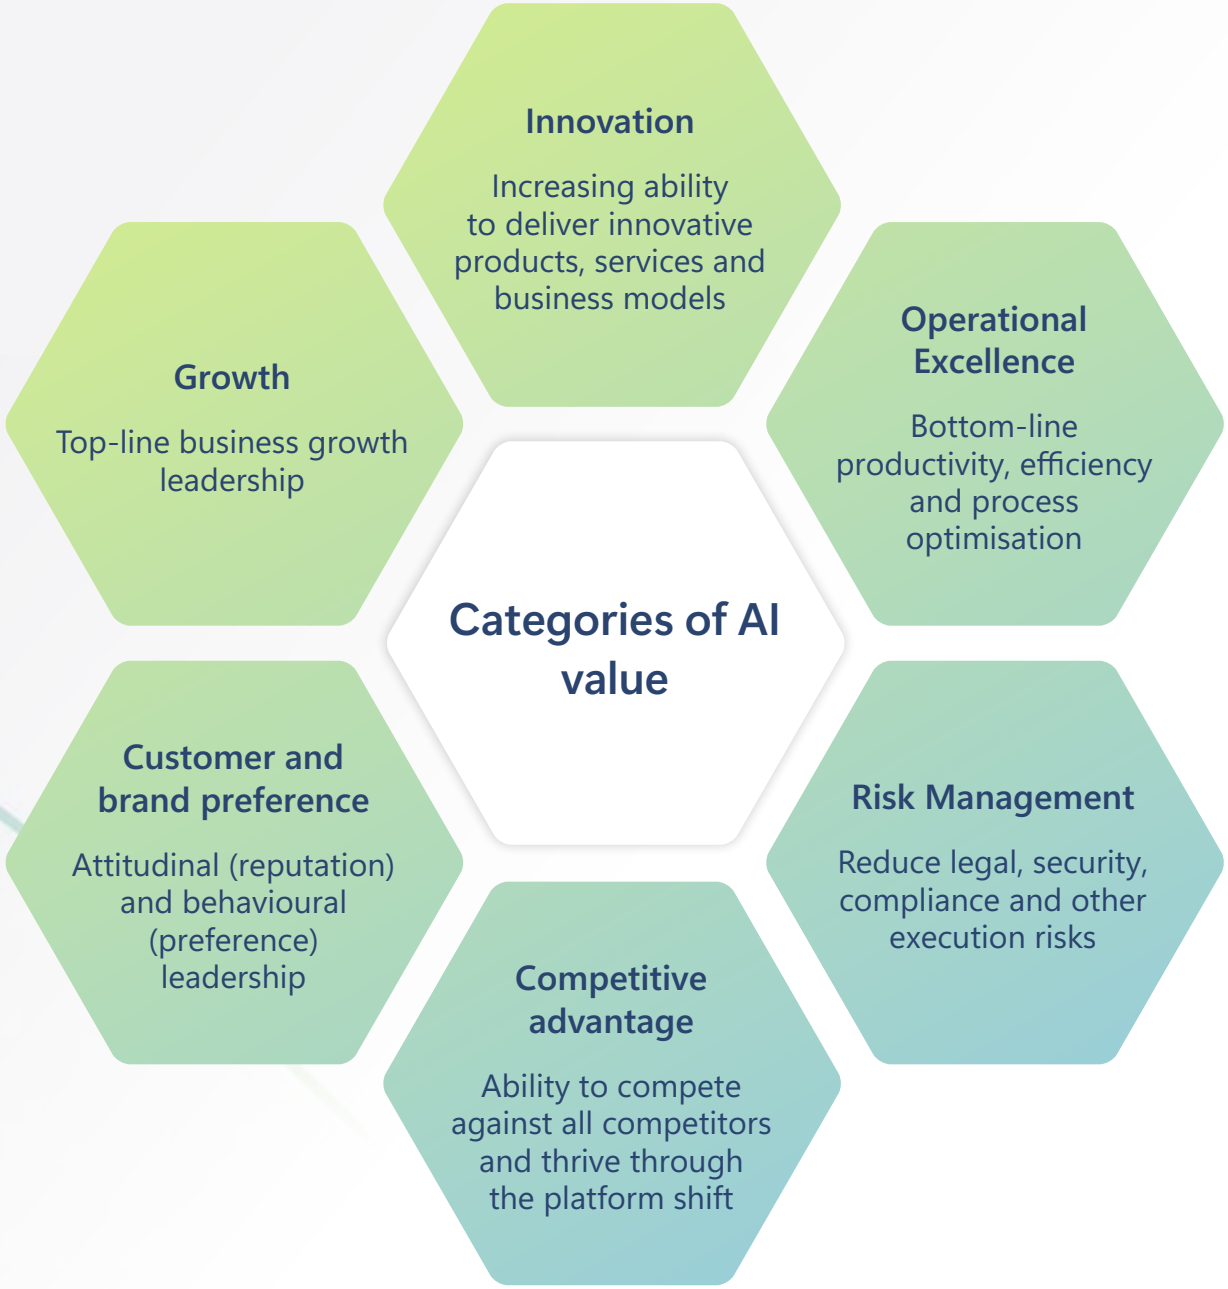

How Azure customers are putting this into practice

Groupama saves time and streamlines work with a generative AI virtual assistant

With 12 million members and €15.9 billion in annual revenue, Groupama Group sought to enhance efficiency in customer interactions as their key ROI metric. Using Azure OpenAI Service and Azure AI Search, the company developed a generative AI-powered virtual assistant to optimise customer relations.

- 80% success rate in delivering verified responses
- Faster, more-effective customer interactions
- Improved employee productivity and job satisfaction

Read the story →

> Our managers save a considerable amount of time and can carry out their work in much better conditions, which enhances the quality of the relationship we maintain with our customers."
>
> **François-Xavier Enderlé**
> Head of Digital Transformation, Groupama

Making this real in your organisation

To unlock the full potential of your AI applications, expand your success criteria beyond direct revenue impact.

Define success beyond revenue

Measure AI's value via efficiency gains, customer satisfaction and innovation-driven outcomes.

Tailor ROI metrics to your industry

Align measurement frameworks with sector-specific AI applications – be it in decision-making, risk reduction or operational efficiency.

Adopt a phased approach

Track early efficiency gains and customer impact while allowing time for broader financial outcomes to develop. Consider long-term impact as well as short-term returns.

Dive deeper: Use this toolkit to assess and define your business needs for building a generative AI app **Making the business case for intelligent applications**.

Insight No. 2

The best use cases for generative AI apps start with a business need

Given seemingly endless use cases, leaders in AI create a process to identify the first one – and often begin with an internal efficiency use case.

Generative AI is an exciting and powerful technology and the pace of innovation continues to accelerate. Nonetheless, the consistent advice from leaders is not to be swayed by excitement. Instead, they suggest focusing first on a clear business need where you can make a measurable impact.

- **For IT leaders,** this means engaging with business stakeholders to understand their most pressing challenges – whether there are inefficiencies in daily workflows (productivity), unmet opportunities (competitive advantage, innovation) or both.
- **For business leaders,** it means working with IT teams to determine which problems generative AI is best suited to solve and which may require a different AI approach – or even a combination of AI techniques.

None of this is easy, even for the most advanced AI experts.

How AI pioneers prioritise use cases

AI leaders adopt a structured approach to pinpointing high-impact AI use cases. They engage stakeholders, assess business needs and align AI investments with strategic goals.

76%

of businesses report needing help identifying the right use case and defining success criteria.[4]

Key takeaways from AI leaders:

- **Prioritise high-impact, low-risk opportunities:** AI pioneers focus first on quick wins – projects that are easier to implement, require fewer resources and deliver measurable benefits early.
- **Engage end users early:** AI pioneers work closely with employees and customers to identify pain points that AI can solve. Understanding real-world challenges helps organisations develop AI applications that deliver immediate value.
- **Use clear criteria to prioritise AI projects:** AI leaders assess potential AI use cases based on five key factors:
  - **Value:** What is the estimated business impact? Could AI increase productivity, reduce risk, improve customer experience or accelerate time to market?
  - **Resources required:** What level of budget, effort and expertise is needed to build and deploy the AI solution?
  - **Specificity:** Is the use case well-defined and practical? The more specific the challenge is, the easier it is to measure success.
  - **Risk level:** Does the use case present significant risks or is it a lower-risk opportunity that can build momentum and organisational confidence?
  - **Opportunity to learn:** Will this project help develop internal AI expertise and prepare the organisation for more ambitious AI initiatives?

"We make a conscious decision to focus on high-impact, low-hanging fruit first. It's all about prioritisation – choosing projects that are easier, quicker and deliver more impact at the start."

Orchestration Leader
Digital AI, Automotive, UK[2]

"For our proof of concept, we focused on real challenges that business users faced rather than theoretical AI capabilities. By addressing specific needs first, we were able to scale AI adoption more effectively."

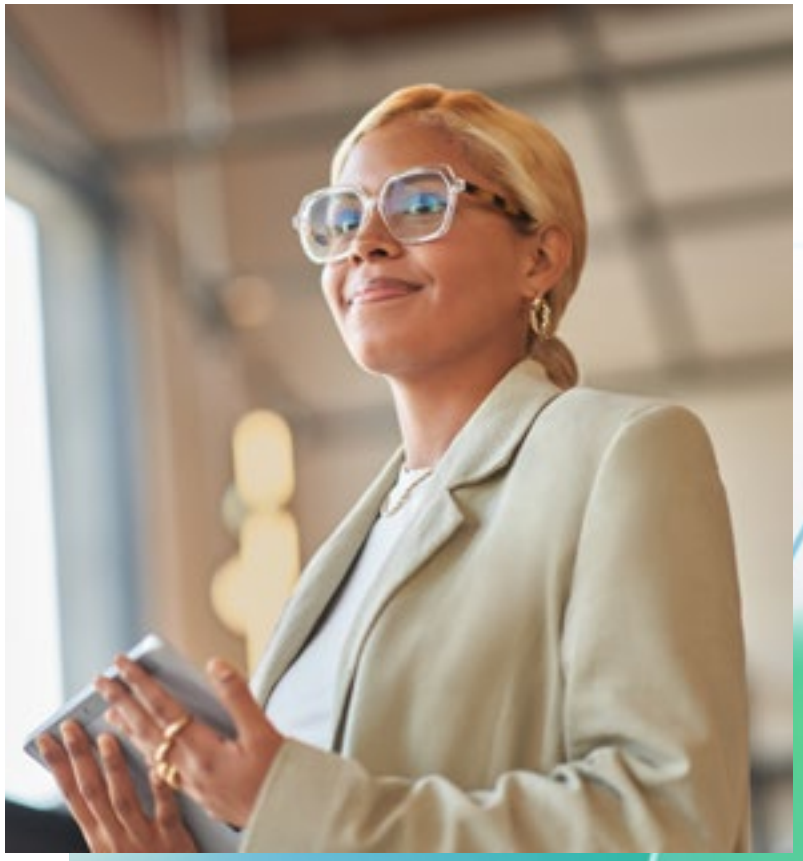

Senior Vice President,
Portfolio Construction & Client Analytics Technology, Financial Services, US[2]

How Azure customers are putting this into practice

Sweco Group leans on experts to rapidly roll out a time-saving AI assistant built in Azure AI Foundry

Sweco, one of Europe's top architecture and engineering firms, wanted rapid deployment for its generative AI solution built with Azure AI Foundry. The team engaged with consultants early on to identify the most valuable automation opportunities for its new digital assistant, SwecoGPT.

- Digital assistant built in just two days
- Improved process efficiency, time savings and error reduction
- Increased productivity from the elimination of mundane tasks

Read the story →

> SwecoGPT is like my assistant. It gives me a head start in writing documents, analysing large piles of documents and even writing emails. It can save me up to two hours or more a day."
>
> **Nena Roes**
> Team Manager of Digital Solutions for Urban Projects, Sweco

Making this real in your organisation

Choosing the right use case is critical to AI success. Organisations should take a strategic approach to identifying AI applications that provide immediate value and align with business priorities.

Conduct surveys to identify issues

Create communication channels where employees and customers can highlight inefficiencies that AI might solve.

Build a backlog of use cases

Keep a log of potential AI applications, even if they're not immediately feasible. As your AI capabilities mature, you can revisit and expand your AI portfolio.

Experiment with AI tools

Use AI trials and proofs of concept to explore different applications and determine which provide the best value.

Dive deeper: Explore how organisations drive impact with AI in **Building AI solutions that drive value**.

Insight No. 3

Generative AI is a team sport between IT and business teams

Bringing multiple stakeholders together accelerates AI success.

Building generative AI applications requires a hybrid approach – blending traditional software development with AI expertise. Even more than in previous tech adoption waves, AI pioneers find that successful implementation hinges on collaboration across business and technical teams.

Even though there's a lot of excitement for the potential of generative AI among businesses, technical leaders at pioneering AI companies warn against falling into the trap of 'if you build it, they will come'. Since people who deeply understand the business and people who deeply understand the technology tend to have different perspectives and areas of expertise, leaders need to bring them together and empower them to drive progress towards a clear goal.

57%

More than half of IT professionals (57%) report having a dedicated team to manage AI.[5]

Organisations that embed cross-functional collaboration into their AI strategies – ensuring that technical and business leaders contribute to development, governance and adoption – are more likely to achieve long-term success. Business leaders help define AI's purpose and ensure its alignment with company goals, while technical teams drive the innovation forwards.

How AI pioneers bridge the gap between business and IT

Leading AI adopters emphasise that cross-functional collaboration accelerates AI success. They have established dedicated teams and governance models to align AI initiatives with business objectives.

Key takeaways from AI leaders:

- **Cross-functional AI councils:** Companies at the forefront of AI adoption create dedicated councils composed of IT, compliance and business leaders. These councils provide strategic oversight, align AI initiatives with business priorities and ensure responsible AI deployment. This can be as scrappy as a 'tiger team' or a formal AI centre of excellence to formalise policies, processes, learning, decision-making and governance across a business unit or enterprise.

  The most commonly cited areas of expertise necessary include:

  - **Senior leadership:** strategic direction, funding and oversight
  - **Business leadership and expertise:** define AI applications that align with company goals and customer needs
  - **Data leadership and expertise:** oversight of and understanding of the organisation's data estate and data strategy
  - **IT and AI:** develop, deploy and maintain AI models, ensuring technical feasibility
  - **Legal and compliance:** establish governance frameworks to ensure responsible data and AI use and ensure that systems meet compliance and regulatory requirements
  - **Security:** understand, plan for, and mitigate potential security risks

- **Dedicated AI governance frameworks:** AI pioneers establish governance policies early to address security, compliance and ethical considerations. These frameworks help organisations manage AI-related risks while maintaining transparency and trust.
- **Embedded legal and compliance teams:** AI leaders integrate legal and compliance stakeholders into the development process from the start. This proactive approach ensures AI applications align with regulatory requirements and company policies, reducing legal and reputational risks.

80%

Through 2027, generative AI will create new roles in software engineering and operations, requiring 80% of the engineering workforce to upskill.[6]

"There is a significant gap between those who understand business needs and those who specialise in AI models. Our role is to bridge this divide, ensuring seamless communication and effective implementation."

Managing Director and Chief Data Architect
Financial Services, US[2]

How Azure customers are putting this into practice

Mars Science & Diagnostics connects experts to build an AI-powered diagnostic tool

With a global shortage of human experts to analyse scans and images, pets had to wait hours and sometimes days before receiving the proper care. Using the Azure AI model catalogue, Azure Kubernetes Service and Azure Cosmos DB – and by bringing together experts from different backgrounds and professionals – Mars Science & Diagnostics built RapidRead to help vets more quickly assess animals with critical conditions.

- 38% average precision boost
- 96.9% precision on a test set
- Accelerated complex workflows

Read the story →

> Working together towards a goal, seeing the alerts for serious findings emerge and routing those directly to a radiologist – it all makes this project incredibly rewarding."
>
> **Michael Fitzke**
> AI Senior Director at Mars Science & Diagnostics

Making this real in your organisation

Successful organisations invest in AI education across teams, ensuring that business leaders understand AI's potential and that technical teams stay aligned with business needs. AI training programmes, internal workshops and AI literacy initiatives help bridge the knowledge gap.

Establish cross-functional teams to accelerate innovation

Rather than centralising all AI work within IT, create cross-functional AI teams that partner closely with business units to ensure AI delivers tangible value across different areas of the organisation.

Create AI committees to vet and prioritise use cases

Bring together stakeholders from legal, HR, compliance and business operations early in the process. These committees review AI use cases, ensuring AI investments align with business goals and regulatory requirements.

Embedded legal and compliance teams

AI leaders integrate legal and compliance stakeholders into the development process from the start. This proactive approach ensures AI applications align with regulatory requirements and company policies, reducing legal and reputational risks.

Dive deeper: Explore strategies for navigating the evolving landscape in The AI Decision Brief.

Insight No. 4

Getting your data house in order is non-negotiable

AI pioneers prioritise mature data management practices as fundamental for success in generative AI

Pioneering AI leaders understand that their success in building successful generative AI applications is contingent on the quality of their data. This data – which could include customer purchase histories, call centre chat logs, research findings, investment strategies and more – is what enables them to build impactful solutions with generative AI.

> Data is where the magic happens. AI models know a lot, but a model doesn't know your company from your competitor until you ground it in your data."
>
> **Jessica Hawk**
> Corporate Vice President, Data, AI and Digital Applications at Microsoft

The criticality of data management cannot be overstated. Before any AI-pioneering organisation unlocked the value of generative AI, they first had to get their data house in order.

How AI pioneers approach data readiness

Leaders have a range of considerations and strategies to prepare their data, but the common theme is clear: the ability to innovate with AI requires data that is accurate, consistent, complete and reliable.

Key takeaways from AI leaders:

- **Data unification is a force multiplier:** unifying your data estate accelerates AI development and unlocks greater business value.
- **A full data audit is required:** in order to improve your overall data quality, conducting a full data audit is essential. During the audit, leaders assess whether their data is consistent, complete and in the right format.
- **Governance must be built in from day one:** legal, compliance, privacy and cybersecurity teams are involved early to ensure customer data is handled properly.

56% In a 2024 survey of IT professionals by Microsoft, 56% stated that data safety and privacy is their top concern with generative AI implementation.[5]

"The first step was massaging our data and building a solid semantic data model. The better your underlying data, the better the answers."

Senior Director
Global Head of Data, Analytics, GenAI & Governance, Professional Services, US[2]

How Azure customers are putting this into practice

Artificial orchestrates labs and augments scientific research with Azure and Azure AI services

Artificial's founders knew they needed to use the cloud to make lab experiments reliable, repeatable and trusted. By migrating its digital lab to Azure, the startup was able to connect its devices, databases and AI tools to deliver better and faster results, using Azure OpenAI Service and Azure AI Search to create digital twins of lab set-ups.

- Avoided expensive wet lab time
- Arrived at more efficient drug candidates
- Enabled real-time oversight of lab environments and experiments

Read the story →

> Using the cloud gives you global visibility and scale, particularly along the data storage front but also with real-time observability… as different scientists manage different parts of the process."
>
> **David Fuller**
> CEO and Co-founder at Artificial

Making this real in your organisation

It's not a new idea, but it's fundamental: the success of any generative AI app is directly tied to the quality and structure of your data. Prioritise getting your data house in order as a prerequisite to AI innovation.

Invest in scalable cloud-based data infrastructure

Since generative AI relies on large amounts of data, centralising your data management in the cloud allows for better scalability and enables multiple teams to access and work with the same data.

Prepare relevant data assets

Build a robust data infrastructure capable of handling large volumes, enabling fast processing and securing access across teams.

Ensure data quality and usability

Conduct a comprehensive data audit to clean, normalise and structure data for use with foundation models.

Involve compliance and governance teams early

Build cross-functional data governance practices from the start to reduce risk and ensure responsible AI use.

Dive deeper: Learn how your proprietary data can drive value and competitive advantage in this on-demand webinar with Accenture: Learn from the Leaders.

Insight No. 5

Testing isn't a "one and done" – and needs to involve end users

Because generative AI is a non-deterministic technology – where the same input doesn't always produce the same output – AI-pioneering leaders adopt an evolved testing approach.

With many other technologies, testing is a phase of the development lifecycle that is mostly completed before the solution reaches the end user. With generative AI, testing happens before, during and after. This has been true of machine learning and neural network systems for years, but with generative AI hitting the mainstream, every organisation needs to grapple with evolving their testing practices.

While the goal is a production-ready solution, achieving that requires extensive testing, feedback and iteration. From hallucination to model drift or mismatches with business context, issues are expected – and solvable. Rather than treating the first versions of a generative AI app as final, successful AI teams treat them as a starting point.

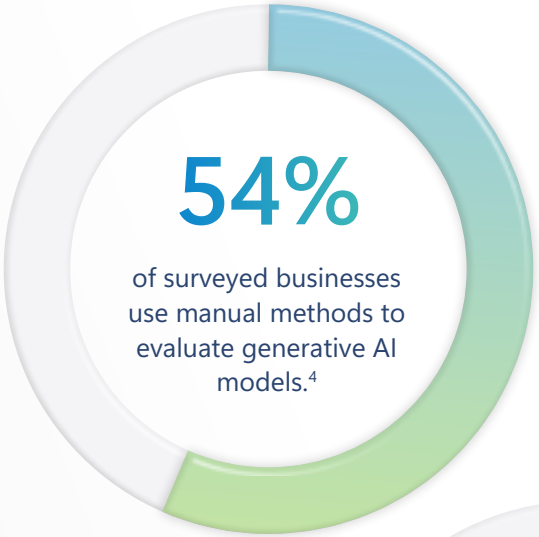

How AI pioneers approach testing

Top-performing AI teams build processes that emphasise continuous learning and iteration. They embrace a "fail fast, learn forward" mindset and build testing loops into every stage of development.

Key takeaways from AI leaders:

- **Start small and learn fast:** Early testing with small groups of pilot users helps identify usability issues, edge cases and unexpected behaviours.
- **Use structured feedback loops:** Successful teams gather input from subject matter experts and business users throughout the pilot to fine-tune model behaviour and align outputs to real-world expectations.
- **Compare against traditional methods:** AI teams conduct parallel testing – comparing AI-generated results to existing workflows – to validate accuracy and highlight areas for improvement.
- **Keep humans in the loop:** While automation is a long-term goal, human input remains vital – especially during pilot phases and ongoing evaluations. Subject-matter experts help validate model performance and flag potential risks.

> "We performed parallel testing by first using our standard method, then applying the same process with the new model. We compared the results to see how the answers differed and evaluated the accuracy."
>
> **Chief of Staff**
> Office of the Chief AI Officer, Financial Services, US[2]

"You need to make sure that you've got an update loop – because you're never finished. It's an ongoing process. We have a data team that constantly checks if we're still getting the right results and whether the model is providing what we need. In the generative AI space, it's not 'develop it and you're done'."

Chief AI Officer
Insurance, UK[2]

How Azure customers are putting this into practice

Daiichi Sankyo uses Azure OpenAI Service to launch an in-house AI system with continuous updates

Global pharmaceutical company Daiichi Sankyo developed a generative AI solution – DS-GAI – to help employees brainstorm, write and improve English-language content. Using Azure OpenAI Service, the team launched a basic interface and added functionality through monthly updates.

- 80% of respondents improved productivity and accuracy
- Updated monthly with new and improved features
- Internal training programmes helped drive adoption

Read the story →

> We do not want to end with the release of DS-GAI, but to make it a driving force for company-wide digital transformation through repeated improvements."
>
> **Mr Yusuke Asoh**
> Associate Director, Daiichi Sankyo Co., Ltd

Making this real in your organisation

Foster a "fail fast, learn forward" culture

Encourage rapid experimentation, knowing that the first version won't be perfect. Position iteration as progress.

Align with existing workflows

Design apps that optimise familiar workflows and ways of working to reduce end-users' learning curves.

Review AI performance regularly

Build recurring review checkpoints to monitor model output, gather user feedback and plan updates.

Dive deeper: Explore strategies for taking AI apps from concept to production in this on-demand webinar.

Conclusion

Launch your generative AI app journey from a scalable, flexible cloud foundation

Success in AI app design goes beyond adopting the leading technology – it requires a holistic approach that includes cultural alignment, cross-functional collaboration, rigorous data preparation and a commitment to continuous iteration. While many of these practices may be familiar, some will guide you into uncharted territory, where the challenges are real and the opportunities for innovation are limitless.

Microsoft Azure is the leading platform for building generative AI apps that deliver real business value. Its powerful AI capabilities, seamless integrations and strong security measures empower you to turn your ideas into solutions that drive competitiveness and long-term growth.

With Azure, you'll find the scalability to handle any workload, flexibility to create customised solutions and robust security to protect your business and data. Plus, with cloud-native tools, you can streamline developer efficiency, enabling faster and easier creation of innovative apps.

Explore Azure's solutions for generative AI app innovation

- **Azure AI Foundry:** Build Copilots and AI solutions faster with prebuilt and customisable models, using your data to innovate at scale.
- **Azure OpenAI Service:** Develop generative AI experiences using prebuilt and curated models from OpenAI, Meta and more.
- **Azure Kubernetes Service:** Develop, deploy and scale cloud-native apps on managed Kubernetes in Azure or at the edge.
- **Azure Cosmos DB:** Power generative AI apps with high-performance and real-time data, ingested and processed at any scale.
- **Azure App Service:** Modernise and scale your applications to accelerate time to market and deliver new experiences.
- **Azure Database for PostgreSQL:** Build AI-ready apps with a fully managed, intelligent and scalable PostgreSQL.
- **GitHub Enterprise:** Seamlessly build, scale, deliver and deploy secure software to the cloud using GitHub's AI-powered developer platform.

Take the next steps

- Find out why Azure is the leading platform for building AI apps
- Read about more innovative use cases for building custom AI applications
- Learn how an AI app goes from idea to reality in our webinar, Taking AI apps from concept to production

Methodology

The quotations in this e-book were drawn from Microsoft customers and leaders, as well as a research study conducted by Emerald Research Group, a subsidiary of the Harris Poll, from 24 October to 22 November 2024, sponsored by Microsoft. The survey sample consists of 29 IT decision-makers, 14 developer decision-makers and 6 business decision-makers from enterprise organisations. These organisations ranged from fewer than 2,500 employees to more than 10,000. The majority of the respondents represent industries such as financial services, technology, manufacturing, insurance, healthcare and retail.

All participants have deployed one or more custom-built generative AI solutions into a production environment and are actively involved throughout the entire AI development process. Participants used a range of generative AI development tools, with no specific toolset required for eligibility. Additional third-party research is cited throughout and appears in the Endnotes section.

References

[1] IDC InfoBrief, sponsored by Microsoft, 2024 Business Opportunity of AI Report, #US52699124, November 2024.

[2] Emerald Research Group: sponsored by Microsoft, GenAI Development Journey, November 2024 This research study was conducted by Emerald Research Group (a subsidiary of the Harris Poll) from 24 October to 22 November 2024, among twenty nine (29) IT decision-makers, fourteen (14) developer decision-makers and six (6) business decision-makers in enterprise organisations. Qualified participants have deployed one or more GenAI custom-built solutions into a production environment and the participants must be actively involved throughout their entire GenAI development process. Participants also had a mix of GenAI development tools and were not required to have used any specific development tools. The opinions and experiences expressed in the online interviews are not representative of the entire market, and only represent those of each individual who participated. All qualitative findings should be interpreted as directional only.

[3] IDC Infographic, sponsored by Microsoft, The Business Opportunity of AI, #US51315823, November 2023.

[4] MIT Technology Review Insights, produced in partnership with Microsoft Azure, Customising generative AI for unique value, 2025.

[5] Microsoft Azure, Tech Pulse: What the rise of AI means for IT professionals, March 2024

[6] Gartner Press Release, Gartner Says Generative AI will Require 80% of Engineering Workforce to Upskill Through 2027, 3 October 2024, https://www.gartner.com/en/newsroom/press-releases/2024-10-03-gartner-says-generative-ai-will-require-80-percent-of-engineering-workforce-to-upskill-through-2027. GARTNER is a registered trademark and service mark of Gartner, Inc. and/or its affiliates in the US and internationally and is used herein with permission. All rights reserved.