

Stanford LLM Lecture Series (watch in order)

- Transformer Fundamentals — <https://lnkd.in/gaaRDexT>
- Transformer Models + Practical Tricks — <https://lnkd.in/gy4FUwNY>
- Transformers → Large Language Models — <https://lnkd.in/gsPiCrEU>
- How LLMs Are Trained — <https://lnkd.in/gvHJvgqP>
- Tuning & Adaptation (fine-tuning, etc.) — <https://lnkd.in/g6kgtPKR>
- Reasoning in LLMs — <https://lnkd.in/gAACSUG6>
- Agentic LLMs (tools, planning, workflows) — <https://lnkd.in/gVm6js9z>
- Evaluation: what "good" really means — <https://lnkd.in/gJhbFQ4s>
- Recap + What's Trending Now — <https://lnkd.in/g5JMNTsf>

Pair the theory with hands-on labs

To make this path more practical and well-rounded, I'd pair the Stanford theory with these free resources:

- Microsoft — Generative AI for Beginners: <https://lnkd.in/dTPBrQb4>
- NVIDIA Developer Program (tools, SDKs, learning): <https://lnkd.in/ejMfduS>
- Google Cloud — ML & AI Training tracks: <https://lnkd.in/eUP8zNik>