

# REGRESSION MODEL CHEAT SHEET

---

## ◆ 1. What is Regression?

Regression is a supervised learning method used to model the relationship between a **dependent variable** and one or more **independent variables**.

---

## ◆ 2. Types of Regression

Type	Description
<b>Linear Regression</b>	Predicts continuous outcome with linear terms
<b>Multiple Regression</b>	More than one predictor variable
<b>Polynomial Regression</b>	Includes non-linear (squared, cubic) terms
<b>Logistic Regression</b>	Classification model using logistic function
<b>Ridge/Lasso</b>	Regularized regression to prevent overfitting
<b>ElasticNet</b>	Combines Ridge and Lasso penalties

---

## ◆ 3. Simple Linear Regression

**Model:**

$$Y = \beta_0 + \beta_1 X + \epsilon$$

**Goal:** Find the line that minimizes residuals (errors).

**Cost Function (MSE):**

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

## ◆ 4. Multiple Linear Regression

**Model:**

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_n X_n + \epsilon$$

## ◆ 5. Key Metrics

Metric	Meaning
$R^2$	% of variance explained by the model
Adjusted $R^2$	Adjusted for number of predictors
RMSE	Root mean square error
MAE	Mean absolute error
p-values	Test significance of coefficients
VIF	Detect multicollinearity

---

## ◆ 6. Assumptions of Linear Regression

1. **Linearity**
  2. **Independence of errors**
  3. **Homoscedasticity (constant variance of residuals)**
  4. **Normality of residuals**
  5. **No multicollinearity**
- 

## ◆ 7. Python Code Snippet

```
python
CopyEdit
import pandas as pd
import numpy as np
import statsmodels.api as sm
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Load Data
df = pd.read_csv("data.csv")

# Define variables
X = df[['X1', 'X2']] # independent variables
y = df['Y']          # dependent variable

# Add constant for statsmodels
X_sm = sm.add_constant(X)

# Fit model (statsmodels)
model = sm.OLS(y, X_sm).fit()
print(model.summary())

# OR using sklearn
lr = LinearRegression()
lr.fit(X, y)
```

```
preds = lr.predict(X)

print("R²:", r2_score(y, preds))
print("RMSE:", np.sqrt(mean_squared_error(y, preds)))
```

---

## ◆ 8. Plotting Diagnostic Charts

```
python
CopyEdit
import matplotlib.pyplot as plt
import seaborn as sns

# Residuals plot
residuals = y - preds
sns.residplot(x=preds, y=residuals, lowess=True)
plt.xlabel("Predicted")
plt.ylabel("Residuals")
plt.title("Residual Plot")
plt.show()
```

---

## ◆ 9. When to Use Regularization

- Many predictors ( $p > n$ )
- High multicollinearity
- Overfitting risk

### Lasso (L1):

$\text{Loss} + \lambda \sum |\beta|$

### Ridge (L2):

$\text{Loss} + \lambda \sum \beta^2$