

IDC MarketScape: Worldwide GenAI Evaluation Technology Products 2025 Vendor Assessment


Nancy Gohring Tim Law

THIS EXCERPT FEATURES IBM AS A LEADER

IDC MARKETScape FIGURE

FIGURE 1

IDC MarketScape Worldwide GenAI Evaluation Technology Products Vendor Assessment



Source: IDC, 2025

Please see the Appendix for detailed methodology, market definition, and scoring criteria.

ABOUT THIS EXCERPT

The content for this excerpt was taken directly from IDC MarketScape: Worldwide GenAI Evaluation Technology Products 2025 Vendor Assessment (Doc # US53015025).

IDC OPINION

This IDC study represents an assessment of the evolving technology used by enterprises to evaluate generative AI (GenAI) applications before moving them into production. While AI evaluation technology is not new, the explosion of interest in GenAI over the past three years spurred the launch of new suppliers and new approaches to GenAI evaluation technology, notably expanding the number and scope of available products. The result is a sector with the characteristics of a very young segment, with wide variations in capabilities, vendor age, revenue, and scope.

Key findings include:

- **The function itself lacks awareness.** Since most enterprises are new to the process of building and deploying GenAI applications, most are unfamiliar with the concept of GenAI evaluation technologies. They face a steep learning curve to understand what GenAI evaluation is, why it's necessary, who should conduct evaluations and when, what techniques are required, and which vendors supply products.
- **There is no common set of capabilities and techniques.** GenAI evaluation products deliver a wide array of capabilities and varying techniques and approaches to supporting GenAI evaluations. Even the language that vendors use to describe similar functionality varies. Over time, a common set of requirements and terminology will likely emerge. But in the meantime, these variations present challenges to enterprises that hope to compare offerings before making a selection.
- **Integrated tools will be required.** As more organizations have begun experimenting with running GenAI evaluations, they've discovered the value of tight integrations with adjacent processes including prompt engineering, offline evaluation, online evaluation, and observability. GenAI evaluation technology products have increasingly expanded into these adjacent sectors or have been added to AI platforms that are used to tune models and build GenAI and AI applications. It is yet unclear exactly which adjacent capabilities will become the most commonly integrated sets offered from a single vendor.

- **The ability to evaluate GenAI apps built using essentially any technology is critical.** Many enterprises will have heterogeneous GenAI applications built using a variety of potential frameworks and development tools, running on a variety of potential cloud or on-premises environments, and using any number of LLMs, all built by multiple suppliers. These enterprises will want evaluation tools that can evaluate GenAI applications regardless of the technology used to build them.
- **Today's GenAI evaluation technology innovations offer insight into future requirements.** A couple of areas stood out as technologies that multiple suppliers were independently developing, and they hint at potential requirements for competitive evaluation tools of the future. One such capability is customizable models that can be used as a judge, with several vendors investing in tools and techniques to aid end users in building such customizations for their particular needs. Another area of technology development is purpose-built databases, with some vendors building their own, designed to scale up to accommodate very large volumes of data, particularly monitoring tools that may be used with evaluation technologies. Another emerging capability for evaluation tools is presenting users with recommendations for fixing problems that evaluation tools surface.
- **Evaluation technology for agentic AI is increasingly important.** Agentic AI has some unique requirements when it comes to running evaluations. For instance, evaluation technology should assess each step of a multi-agentic workflow, as well as the final output, to determine where problem areas might be occurring. Evaluation data set requirements are more complex in agentic workflows, and as such, synthetic data generation becomes more important. Many of the vendors in this document are offering these and other techniques, with many of the capabilities introduced in CY25.

IDC MARKETSCAPE VENDOR INCLUSION CRITERIA

To determine the set of suppliers included in this research, IDC identified vendors of evaluation technologies that met the following criteria:

- Have 2024 annual revenue of \$1+ million
- Offers a product for evaluating GenAI apps in development (offline evaluations)
- Offering predominantly based on company IP (as opposed to open source or other software developed externally)
- Product commercially available during CY24

ADVICE FOR TECHNOLOGY BUYERS

Recognizing how important evaluation is to the potential success of a GenAI or agentic AI rollout is critical to the success of these AI applications in the enterprise. Testing applications with a small set of prompts developed via guesswork will not deliver the level of quality and accuracy that enterprises require from GenAI and agentic AI applications. Organizations that recognize that evaluation is a key component of the GenAI development and operating workflow have a better chance at developing and supporting applications that deliver on target goals.

When selecting an evaluation technology that will best serve your organization, consider the following:

- **Look at how heterogeneous your GenAI environment is and how heterogeneous you expect it will be in the future.** Some evaluation products are most useful in evaluating applications that were built using tools supplied by the same vendor offering the evaluation product. Organizations that expect to stick with that vendor's environment for the long term may be well suited to adopting these evaluation products, particularly if they're well integrated with the rest of the supplier's suite of offerings. If you expect your applications to be built on technologies from a wide variety of suppliers, consider vendors that emphasize their neutrality and prioritize out-of-the-box integrations that support a variety of technologies to avoid a high volume of manual integration work.
- **Evaluate the level of internal expertise you have, and match that to the offering.** Some evaluation offerings are targeted at the many organizations that are new to evaluations, with features that are designed to be easy to use and to guide users toward best practices. Others are barely more than a set of SDKs, which may suit experienced developers.
- **Consider the supplier's ability to be nimble and innovative.** As with practically all aspects of GenAI, evaluation technologies is a fast-evolving segment. Many vendors have unique approaches and techniques and have moved quickly to introduce features targeted at agentic AI use cases. The vendors that quickly support new technology developments and that deliver innovative capabilities will be best suited to supporting your evolving needs into the future.
- **Identify the personas for whom you want to use the evaluation tools and whether the products you're assessing cater to them.** GenAI evaluation products should be used by the developers of GenAI applications, but they also can be useful to business personas who have deep understanding of use cases. For enterprises that recognize the value of a tight collaboration between technical and business teams in the development of GenAI applications, examine

the features of evaluation tools that are targeted at the business persona's interest.

VENDOR SUMMARY PROFILES

This section briefly explains IDC's key observations resulting in a vendor's position in the IDC MarketScape. While every vendor is evaluated against each of the criteria outlined in the Appendix, the description here provides a summary of each vendor's strengths and challenges.

IBM

After a thorough evaluation of IBM's strategies and capabilities, IDC has positioned the company in the Leaders category in this 2025 IDC MarketScape on GenAI evaluation technology products.

IBM's evaluation technology is available as part of watsonx.governance and it is embedded in IBM's concept of the GenAI and agentic AI development life cycle tightly integrated with adjacent IBM capabilities.

As an example that demonstrates the value of IBM's life-cycle perspective, a business leader at an IBM customer can start the process of developing a new GenAI application in watsonx.governance by defining a business problem and uncovering potential risks using a guided questionnaire available in the product. The risk dimensions are translated into metrics that can be used during the evaluation process to measure the identified potential risks and recommend mitigations. As those metrics are collected throughout the evaluation process, metadata is extracted and stored in a fact sheet, which serves as a record to track the life cycle of the application development process, capturing information about the model, prompt templates, and so on. This approach of offering evaluation technology as part of a tightly integrated life cycle that includes documentation is a differentiator that aims at enterprise users.

watsonx.governance comes with a long list of metrics available out of the box, including drift identification, and customers can build custom metrics. The Evaluation Studio guides users through a process of developing custom rankings, where the user can select metrics and assign weights, and then the software creates a custom formula for applying the custom ranking. IBM has trained small models that serve as LLM judges and supports the creation of custom judges. Users can bring their own models to serve as judge as well.

To help end users identify the source of poor performance discovered during an evaluation, watsonx.governance offers a Venn diagram to visualize the intersection of multiple factors that may contribute to the poor performance. The Venn diagram

shows correlated metrics and then surfaces the data with poor results so that users can drill into the problematic data to determine the cause of the problem. For evaluating security weaknesses, the offering surfaces mitigation recommendations.

Evaluation Studio offers experiment management and a UI for comparing prompt templates, and users can manage data sets from within watsonx.governance. Customers that use watsonx.governance guardrails can collect data that guardrails block to add it to experiment data sets. Some dashboards come out of the box and users can also create their own.

Earlier this year, IBM added capabilities for evaluating agents, including those in production. One technique, called Agent Governance In the Loop Tool, evaluates agent tool usage, and customers can use the resulting metrics to guide the agent's workflow. For instance, if the tool determines that context relevance is low, the agent can be instructed not to trust the context and instead do a web search, for example. IBM also released agent evaluators, including root cause analysis and plans, to roll out experiment tracking for agents and agentic AI production monitoring.

Strengths

IBM has a long history as an AI innovator and an enterprise supplier and as such has a feature-rich offering. It offers a few differentiators, including root cause analysis and automated features like the fact sheet that appeals to enterprise use cases.

Challenges

IBM has a large catalog of offerings and it's not always clear which can be used independently and then with which features. While deep integration with adjacent capabilities is a plus for some users, others that might only want the evaluation technologies may be uncertain about their options.

Consider IBM When

Enterprises that have a diverse technology environment may find that IBM represents a neutral supplier — it's not tied to a particular cloud service, for instance. Also, enterprises that value the broader set of adjacent IBM offerings, including the automated documentation, guardrails, and security offerings, should consider IBM.

APPENDIX

Reading an IDC MarketScape Graph

For the purposes of this analysis, IDC divided potential key measures for success into two primary categories: capabilities and strategies.

Positioning on the y-axis reflects the vendor's current capabilities and menu of services and how well aligned the vendor is to customer needs. The capabilities category focuses on the capabilities of the company and product today, here and now. Under this category, IDC analysts will look at how well a vendor is building/delivering capabilities that enable it to execute its chosen strategy in the market.

Positioning on the x-axis, or strategies axis, indicates how well the vendor's future strategy aligns with what customers will require in three to five years. The strategies category focuses on high-level decisions and underlying assumptions about offerings, customer segments, and business and go-to-market plans for the next three to five years.

The size of the individual vendor markers in the IDC MarketScape represents the market share of each individual vendor within the specific market segment being assessed.

IDC MarketScape Methodology

IDC MarketScape criteria selection, weightings, and vendor scores represent well-researched IDC judgment about the market and specific vendors. IDC analysts tailor the range of standard characteristics by which vendors are measured through structured discussions, surveys, and interviews with market leaders, participants, and end users. Market weightings are based on user interviews, buyer surveys, and the input of IDC experts in each market. IDC analysts base individual vendor scores, and ultimately vendor positions on the IDC MarketScape, on detailed surveys and interviews with the vendors, publicly available information, and end-user experiences in an effort to provide an accurate and consistent assessment of each vendor's characteristics, behavior, and capability.

Market Definition

GenAI evaluation technology is a set of tools and methods used to assess the performance, accuracy, fairness, and other qualities of artificial intelligence systems. This technology is designed to help developers of GenAI applications understand how well their models are working, identify areas for improvement, and ensure that the model is aligned with their goals and values. GenAI evaluation technology can include tools for testing and debugging GenAI models, as well as techniques for measuring and analyzing the output of GenAI applications.

Evaluation technologies are increasingly important in delivering confidence to enterprises that are reluctant to deploy GenAI applications because of concerns about risk related to accuracy, hallucination, privacy, and security. Evaluation technologies support engineers in determining whether changes should be made to the application,

including selecting a different model or prompt engineering, to deliver better outcomes. IDC anticipates that organizations will update some applications on a regular basis, requiring regular use of evaluation technologies to ensure that such updates continue to deliver desired outcomes.

LEARN MORE

Related Research

- *Global GenAI Technology Trends Survey, May 2025* (IDC #US53177225, July 2025)
- *IDC Quick Poll, 2Q25: Agentic AI and Data* (IDC #US53638725, June 2025)
- *IDC Market Glance: GenAI Monitoring, 2Q25* (IDC #US53359125, May 2025)
- *Agentic AI Impact on Enterprises: From the Tech Stack to the Future of Work and Services* (IDC #US53272524, April 2025)
- *IDC MaturityScape Benchmark: AI-Fueled Organization Worldwide, 2025* (IDC #US53271725, March 2025)

Synopsis

This IDC study assesses the GenAI evaluation technology products market, an evolving segment that has emerged in support of the growing enterprise interest in GenAI applications. It examines 13 vendors, highlighting technical capabilities as well as challenges and potential.

"Evaluations are key to the GenAI and agentic AI application development process, allowing developers to improve accuracy, reduce hallucinations and, ultimately, achieve business objectives," said Nancy Gohring, senior research director, AI, IDC. "The evaluation tool supplier segment is evolving quickly to meet this need, with new entrants and established AI development platforms competing for attention in this low-visibility sector."

ABOUT IDC

International Data Corporation (IDC) is the premier global provider of market intelligence, advisory services, and events for the information technology, telecommunications, and consumer technology markets. With more than 1,300 analysts worldwide, IDC offers global, regional, and local expertise on technology, IT benchmarking and sourcing, and industry opportunities and trends in over 110 countries. IDC's analysis and insight helps IT professionals, business executives, and the investment community to make fact-based technology decisions and to achieve their key business objectives. Founded in 1964, IDC is a wholly owned subsidiary of International Data Group (IDG, Inc.).

Global Headquarters

140 Kendrick Street
Building B
Needham, MA 02494
USA
508.872.8200
Twitter: @IDC
blogs.idc.com
www.idc.com

Copyright and Trademark Notice

This IDC research document was published as part of an IDC continuous intelligence service, providing written research, analyst interactions, and web conference and conference event proceedings. Visit www.idc.com to learn more about IDC subscription and consulting services. To view a list of IDC offices worldwide, visit www.idc.com/about/worldwideoffices. Please contact IDC at customerservice@idc.com for information on additional copies, web rights, or applying the price of this document toward the purchase of an IDC service.

Copyright 2025 IDC. Reproduction is forbidden unless authorized. All rights reserved.